visualizemi: Visualization, Effect Size, and Replication of Measurement Invariance for Registered Reports

Erin M. Buchanan[1]

[1] Harrisburg University of Science and Technology

Author Note

Abstract

Latent variable modeling as a lens for psychometric theory is a popular tool for social scientists to examine measurement of constructs (Beaujean, 2014). Journals such as *Assessment* regularly publish articles supporting measures of latent constructs wherein a measurement model is established. Confirmatory factor analysis can be used to investigate the replicability and generalizability of the measurement model in new samples, while multigroup confirmatory factor analysis is used to examine the measurement model across groups within samples (Brown, 2015). With the rise of the replication crisis and "psychology's renaissance" (Nelson et al., 2018), interest in divergence in measurement has increased, often focused on small parameter differences within the latent model. This manuscript presents `visualizemi`, an R package that provides functionality to calculate multigroup models, partial invariance, visualizations for (non)-invariance, effect sizes for models and parameters, and potential replication rates compared to random models. Readers will learn how to interpret the impact and size of the proposed non-invariance in models with a focus on potential replication and how to plan for registered reports.

Keywords: multigroup confirmatory factor analysis, measurement invariance, visualization, effect size

visualizemi: Visualization, Effect Size, and Replication of Measurement

Invariance for Registered Reports

Psychological assessments play a critical role in our ability to measure and analyze constructs to support theories and experimental hypotheses. Defining and creating assessments to validly and reliability measure constructs is often difficult because phenomenon, such as anxiety, are often not directly observable. Instead, we use surveys and questionnaires to indirectly assess the underlying construct (DeVellis & Thorpe, 2022). Latent variable modeling (i.e., structural equation modeling) is a popular tool for the validation of developed survey instruments to verify scale dimensionality, structure, and model fit. A simple search for scale development reveals thousands of articles in psychology that examine new and previously published work, thus, illustrating the interest in both measurement and the use of validation techniques. Unfortunately, except in specialty journals, much of the validity evidence and/or development for measures used in empirical studies is not reported within the journal article (Barry et al., 2014; Weidman et al., 2017). Without this information, it is difficult to interpret individual study conclusions, as validity information allows for judgment of usefulness of the measured values (Flake & Fried, 2020). Further, the current focus on replication (Makel et al., 2012; Makel & Plucker, 2014; Zwaan et al., 2018), reproducibility (Nelson et al., 2018), and the credibility of our results (Vazire et al., 2022) has demonstrated questionable measurement practices - decisions that researchers make like survey selection and scoring that impact the results of the study (Flake & Fried, 2020). Transparent reporting of the use and creation of scales can improve both interpretation and reproducibility when using surveys developed to measure latent constructs (Shadish et al., 2001).

A secondary concern for developed measures is the potential for differential responding and assessment within target populations. For example, Trent et al. (2013) examined for potential variability in the Revised Child Anxiety and Depression Scale in

54 White and Black youths (Chorpita et al., 2000). They found that the scale mostly

55 functioned the same for both White and Black individuals but differences in averages on

56 individual items could potentially affect the scoring and interpretation of the scale results.

57 This comparison of sub-populations is the test of measurement invariance (Meredith, 1993).

58 Invariance or equivalence implies that the scale operates in the same fashion for each

59 sub-group, and thus, differences in the final latent variable scores can interpreted as

60 differences in populations. Non-invariance suggests that individuals respond or interpret

61 items differently, and thus, differences in scores may represent different scores on the latent

62 variable in the population or differences in measurement. Non-invariant measurement may

63 lead to misleading results when making group comparisons, and assessing invariance has

64 become a popular technique in scale development (Van De Schoot et al., 2015).

65   Measurement invariance has been explored and implemented for the last fifty years

66 (Jöreskog, 1971; Sörbom, 1978) and implemented in the most popular structural equation

67 modeling programs (Boker et al., 2011; Jöreskog & Sörbom, 2001; Rosseel, 2012). Byrne et

68 al. (1989) extended the ideas of multigroup testing by suggesting partial invariance (followed

69 by Meredith, 1993). Partial invariance occurs when non-invariance is found but can be

70 attributed to only a few parameter estimate differences between groups (i.e., items 1 and 2

71 have different factor loadings but all others are the same). This testing provided an

72 advantage to understand where the potential non-invariance may occur for further study and

73 interpretation guidelines. As the field pushes back against favoring cutoff criteria and rules

74 of thumb (Marsh et al., 2004; Putnick & Bornstein, 2016), an effect size measure for

75 translating "how much" non-invariance was developed d_{MACS} (Nye & Drasgow, 2011). This

76 effect size examines the differences in observed variables between the two groups for both the

77 factor loading and the item intercept; thus, any differences in either or both will increase the

78 effect size for non-invariance (Stark et al., 2006).

79   With d_{MACS} and measurement invariance testing, researchers can begin to quantify

80 how and where their construct measurement may vary between groups. Yet, given the large

81 number of studies that show non-invariance, it is clear that equivalence can be hard to meet.

82 It is difficult to know if non-invariance occurs because of random sampling error, true

83 population differences, or differences in replication and reproducibility of the construct in a

84 new sample. The field of psychology is increasingly interested in pre-registration (i.e.,

85 registering plans for a study before data collection, Nosek et al., 2018) and the promotion of

86 transparency in study design, implementation, and analysis (Mayo-Wilson et al., 2021), in

87 addition to supporting replication studies (Zwaan et al., 2018). Registered (replication)

88 reports provide an advantageous avenue for the pre-registration of measurement tests, as

89 they allow a researcher the ability to have their study accepted in principle, regardless of the

90 results of a test of construct validity, reliability, or measurement invariance (Hobson, 2019;

91 Nosek & Lakens, 2014). However, there are few tools that can provide effect size measures

92 for models, individual parameters, or visualization for researchers to plan for future studies.

93 d_{MACS} provides the opportunity to begin to think about the smallest effect size of interest or

94 the smallest meaningful effect size for measurement invariance and replication (Anvari &

95 Lakens, 2021; i.e., two studies with overlapping confidence intervals "replicate," even if the

96 test of measurement invariance does not, Lakens, 2017). As mentioned, d_{MACS} has only

97 really been explored for a combined intercept and loadings, and while useful, does not

98 necessarily allow a researcher to pinpoint specific issues within an observed variable.

99 Therefore, purpose of this manuscript is to describe an R package, `visualizemi`, that

100 provides functionality to calculate multigroup confirmatory factor analysis, partial invariance

101 tests, visualizations of the size of non-invariance, and potential effect sizes for overall models

102 and individual parameters. No known visualization techniques have been proposed for

103 measurement invariance. By creating panel visualizations, we can supplement a researcher's

104 ability to judge the strength of the non-invariance differences and effect size for each item.

105 The proposed effect sizes demonstrate the likelihood of replication with a similar sample as

106 compared to a randomly assigned group model, thus, illustrating what type of measurement

one might expect to find, and how different that is from random chance. Within this technique, the individual parameter effect sizes can be calculated: both the group differences within a model as compared to random and the likelihood of a parameter replication compared to random groups. Coupled with other indicators (i.e., fit indices differences, d_{MACS}), we can move toward a better understanding of how much measurement non-invariance is meaningful. This tutorial and package will help researchers plan future studies and aid in the ability to estimate a smallest effect of interest for measurement invariance studies, rather than relying on fit indices and rules of thumb alone.

By the end of this tutorial manuscript, readers will:

1. Learn how to use *visualizemi* to analyze multigroup confirmatory factor analysis, examine partial for invariance, and create visualizations of parameters.
2. Learn how to estimate the potential replication of multigroup models and their parameters using bootstrapping compared to a random group model.
3. Be able to calculate and interpret effect sizes for model and parameter replication, as well as parameter group differences.
4. Understand the impact of measurement variability on replication and generalizability.

The tutorial will start with a general overview of relevant topics to orient readers to invariance testing and MGCFA effect sizes. Next, the reader will learn about the `visualizemi` package functions for 1) running the multigroup analysis, 2) running a partial invariance analysis, 3) plotting the partial invariance, 4) estimating replication and effect sizes at the model level, and 5) estimating replication and effect sizes at the parameter level. Last, data from Aiena et al. (2014) and Chen et al. (2020) examining the measurement invariance of the RS-14 (Wagnild, 2009) will be used to demonstrate the application of the package on real data. The *visualizemi* package vignette includes an additional tutorial walk through.

₁₃₂ **Terminology**

₁₃₃ **MGCFA**

₁₃₄   Multigroup confirmatory factor analysis (MGCFA) was proposed as a method to

₁₃₅ examine differences in scale functioning across groups (Brown, 2015) using structural

₁₃₆ equation modeling and confirmatory factor analysis (CFA). The goal of MGCFA is to

₁₃₇ determine if groups are invariant or show the same response patterns on the scale. MGCFA

₁₃₈ is often performed in sequential steps to determine the location and impact of potential

₁₃₉ non-invariance. The most common procedure starts by examining overall scale structure for

₁₄₀ each group individually to show that the groups can be combined into one nested model (i.e.,

₁₄₁ both models at least converge). The combined model, often described as the *configural*

₁₄₂ model, creates one CFA for both groups that allows each group's estimated parameters to

₁₄₃ vary. Equality constraints between group parameters are then added sequentially to the

₁₄₄ model (Brown, 2015). If parameters are found to be equivalent between groups, these models

₁₄₅ are considered "invariant", suggesting that any subtle differences in the parameter estimate

₁₄₆ should not effect overall scoring and assessment.

₁₄₇   The first equality constraint added is usually the item factor loadings (*metric* model)

₁₄₈ which requires items to be related to the latent variable at the same strength across groups.

₁₄₉ The intercepts are then constrained to be equal across groups (*scalar* model) to determine

₁₅₀ the invariance of item averages. The item variances (*strict* model) can then be examined to

₁₅₁ determine if the general variation in item answers are equal across groups. Other parameter

₁₅₂ equality constraints can be set at the latent variable level (latent means, variance,

₁₅₃ covariances), but the focus is generally on the observable variables and their parameter

₁₅₄ estimations. Non-invariance is determined by examining differences in fit indices between

₁₅₅ models with constraints versus without (Cheung & Rensvold, 2002).

Partial Invariance

If a specific step within the MGCFA testing framework shows non-invariance, partial invariance is often used to investigate where and "how much" invariance occurs (Byrne et al., 1989; Meredith, 1993). Similar to *post hoc* follow up tests in ANOVA, each equality constraint for that model is examined one at a time by allowing groups to vary. If the model fit improves when groups are allowed to have separate parameter estimations, the item parameter estimate is considered non-invariant. When only a few parameters are found that impact invariance, models are considered partially invariant, implying that most, but not all parameters are equal across groups. The researcher then interprets the impact of those items and parameters on overall scores and assessment usefulness. Partial invariance investigation is a useful tool for finding specific items that vary between groups, but does not fully explain the effect size of the difference between groups and the impact on the overall model.

Effect Size: d_{MACS}

d_{MACS} was developed from Differential Item Functioning (DIF) measurement in Item Response Theory (IRT, Stark et al., 2004) wherein the effect size (DIF) portrays the group differences that lead to overall item score differences. Item scores can be mathematically defined as:

$$\hat{X}_{iR} = \tau_{iR} + \lambda_{iR}\xi$$

i indicates an individual item, and R indicates the reference group (Nye & Drasgow, 2011). Therefore, an individual score is predicted by the intercept of the item (τ) added to the item loading (λ) times the latent variable score (ξ). Nye and Drasgow (2011) demonstrated that the terminology from IRT and DIF can be used to create a measure of item functioning differences within CFA, d_{MACS}. d_{MACS} is then the difference of reference group versus focal group (e.g., group 1 versus group 2) divided by a pooled standard deviation similar to formulas for d proposed by Cohen (2013). Therefore, if τ or λ are

180 individually (or both) different across groups, it will impact all the predicted X_{iR} scores, and

181 thus, impact d_{MACS} values. While d_{MACS} represents a necessary step for development of

182 effect sizes within CFA, it does not separate the differences in parameters between groups in

183 a way that can be paired with traditional MGCFA testing and partial invariance. In

184 addition to the missing effect sizes at model and parameter levels, no effect size to date gives

185 the researcher a feeling for potential replication of the invariant or non-invariant items. Last,

186 in line with general suggestions by Cumming (2012) and Cumming and Calin-Jageman

187 (2016), the visualization of effect sizes in MGCFA would be an added tool for researchers to

188 gauge the size of group differences.

Package Functions

190   The complete code for this manuscript can be found at https://osf.io/wev5f/. This

191 tutorial was registered at https://osf.io/vwf4d, and the example provided at the end of the

192 manuscript was added after that registration. The R package and replication/effect sizes was

193 added after the original manuscript submission. The simulation study used to design

194 plotting functions and test effect sizes can be found in the supplemental materials, along

195 with worked code examples.

MGCFA: `mgcfa()` Function

197   First, we would create our model code in *lavaan* syntax (Rosseel, 2012). The

198 *visualizemi* package does generally require raw data for bootstrapping purposes, and an

199 example of how to simulate data from models and covariance/correlations tables that

200 sometimes are provided in manuscripts (rather than the raw data) is provided in the

201 supplemental documentation. The `mgcfa()` function is designed to flexibly allow you to

202 leverage *lavaan*'s package functions to calculate multiple measurement steps at once. You

203 would include:

204   1) the model syntax in the `model` argument.

205   2) the dataframe in the `data` argument of our functionl

3) the name of the grouping variable in quotes for `group`.

4) and the equality constraints you would like to impose in order in `group.equal`.

5) `...` any other *lavaan* arguments you would like to use such as `meanstructure` or `estimator`.

The following output is saved:

1) `model_coef`: The parameter estimates for each model with the model step included in a *model* column. This set of coefficients can be used for other functions. This dataframe is created with *broom*'s `tidy()` function if you wish to recreate this table without running the `mgcfa()` function (Robinson et al., 2023).

2) `model_fit`: The model fit indices from `fitmeasures()` to review for overall model fit and invariance judgments. The name of the model is included in a *model* column.

3) `model_overall`: A saved *lavaan* fitted model of all groups together without any equality constraints or grouping variables. These objects can be used with any function that normally takes a saved model: `parameterEstimates()`, `modificationIndices()`, `semPlot::semPaths()`, and so on (Epskamp, 2022).

4) `group_models`: A list of saved fitted models for each group separately.

5) `model_configural`: A saved fitted model for the configural model that nests together each group into one model with no other constraints.

6) `invariance_models`: A list of saved fitted models that consecutively adds `group.equal` constraints.

Partial Invariance: `partial_mi()` Function

The `partial_mi()` function aids in the calculation of partial invariance for a specific step of the MGCFA process. The function includes the following arguments:

1) `saved_model`: The saved *lavaan* model with the equality constraints at the level of measurement invariance you would like to examine for partial invariance.

231   2) `data`: The dataframe where the model was estimated.

232   3) `model`: The model syntax for the overall model.

233   4) `group`: The grouping variable column in the dataframe.

234   5) `group.equal`: The equality constraints including in your original multigroup tests.

235   6) `partial_step`: The level of partial invariance you wish to test.

236   In this function, each parameter with the appropriate *lavaan* syntax is relaxed

237   individually (i.e., `~1` for intercepts, `~~` for residuals, etc.). The fitted models are saved in the

238   `models` output, and the `fit_table` output includes all fit indices for each model to

239   investigate potential areas of partial invariance based on the researcher's desired criterion.

240   **Visualization of Invariance: `plot_mi()` Function**

241   Once we know which items are non-invariant, the `model_coef` output from the

242   `mgcfa()` can be used directly in `plot_mi()`. The plot outputs will be described below. First,

243   here are the arguments for the function:

244   1) `data_coef`: A tidy dataframe of the parameter estimates from the models. This

245   function assumes you have used `broom::tidy()` on the saved model from *lavaan* and

246   added a column called "model" with the name of the model step (Robinson et al.,

247   2023). This function will only run for models that have used the grouping function (i.e.,

248   configural, metric, scalar, and strict or other combinations/steps you wish to examine).

249   2) `model_step`: Which model do you want to plot? You should match this name to the

250   one you want to extract from your model column in the `data_coef`.

251   3) `item_name`: Which observed variable from your model syntax do you want to plot?

252   Please list this variable name exactly how it appears in the model.

253   4) `x_limits`: What do you want the x-axis limits to be for your invariance plot? The

254   default option is to assume the latent variable is standardized, and therefore, -1 to 1 is

255   recommended. Use only two numbers, a lower and upper limit. This value also

256   constrains the latent mean diagram to help zoom in on group differences because the

257 scale of latent means is usually centered over zero. You can use this parameter to zoom

258 out to a more traditional histogram using `c(-2, 2)`.

259 5) `y_limits`: What do you want the y-axis limits to be for your invariance plot? Given

260 that the latent variable is used to predict the observed values in the data, you could

261 use the minimum and maximum values found in the data. If that range is large,

262 consider reducing this value to be able to visualize the results (i.e., otherwise it may be

263 too zoomed out to judge group differences). Use only two numbers, a lower and upper

264 limit.

265 6) `conf.level`: What confidence limit do you want to plot? Use $1 - \alpha$.

266 7) `model_results`: In this argument, include the saved *lavaan* output for the model

267 listed in the `model_step` argument.

268 8) `lv_name`: Include the name of the latent variable, exactly how it is listed in your

269 *lavaan* syntax. You should plot the latent variable that the `item_name` is linked to. If

270 you have items that load onto multiple latent variables, you will need to make multiple

271 plots.

272 9) `plot_groups`: If you include more than two groups in a multigroup model, the

273 automatic assumption is that you want the first two groups for this visualization. If

274 not, include the names of the groups here to plot.

275 The outputs from this function are several *ggplot2* objects that can be edited or saved

276 directly using *ggplot2* functionality (Wickham, 2016).

277 1) `complete`: The output from this model can be found in Figure 1. On the left-hand side,

278 the item invariance is plotted, and on the right-hand side, the latent mean distributions

279 for the two groups are plotted. In the item invariance sub-plot, the visualization

280 includes all three components traditionally seen in MGCFA testing steps: loadings,

281 intercepts, and residuals. Each visualization element was designed to match the

282 traditional visualization for that type of output. All parameter estimates are plotted

283   on the unstandardized estimates and their confidence interval based on the standard

284   error of the estimate. All plots are made with *ggplot2* and *cowplot* (Wilke, 2020).

Figure 1

Invariant model visualization demonstrating the components of the `plot_mi()` *function in* visualizemi.

285   2) `intercept`: Only the left-hand side of the complete plot designed to represent

286   intercepts and factor loadings. Factor loadings represent the slope of the regression

287   equation for the latent variable predicting the scores on the observed variable

288          $(\hat{Y} \sim b_0 + b_1 X + \epsilon)$. The y-axis indicates the observed variable scores, and here, the

289   plot includes the entire range of the scale of the data this simulated item. The ribbon

290   bands around the plotted slopes indicate the confidence interval for that estimate. In

291   this plot, while the coefficients for each group are not literally equal, the overlapping

292   and parallel slope bands indicate they are not different practically.

293   The item intercepts (b_0) are plotted on the middle line where they would cross the

294   y-axis at a latent variable score of zero. These are represented by a dot with a set of

295   confidence error bars around the point. In this invariant depiction, the overlap in the

296   intercepts is clear, indicating they are not different. You can use `y_limits` to zoom in on

297   the graph if these are too small to be distinguishable. 3) `mean`: The right-hand side of the

298   complete plot graphing the latent variable means and density from the data. The latent

variable is shown on the x-axis using standardized values (i.e., *z*-scores) where -1 indicates one standard deviation below the mean for the latent variable, 0 indicates the mean for the latent variable and so on. The lines indicate the means of the latent variables from the simulated dataset. Group labels are represented in the figure caption on the bottom. Group 1 is usually the group that is alphabetically first in the data set or whichever group is the first that appears when using the `levels()` command. 4) `variance`: A split geom violin plot indicating the variance distribution of the plotted item. Residuals are trickier to plot, as they are the left over error when predicting the observed variables ϵ. It is tempting to plot this value as the confidence band around the slope, however, that defeats the purpose of understanding that the slopes are estimated separately from the residuals, and both have an associated variability around their parameter estimate. Therefore, residuals are represented in the inset picture at the bottom right of the item invariance plot. The black bars represent the estimated residual for each group. The distributions are plotted to represent the normal spread of values using the standard error of the residuals. The violin plot allows for direct comparison of those residuals and their potential distributions. Note that the placement has nothing to do with the x or y-axis and is designed to always show in the same location, regardless of size/value. The plots are included separately so they can be arranged in a different fashion if desired.

Model Replication and Effect Sizes: `bootstrap_model()` Function

The `bootstrap_model` function in *visualizemi* was designed to estimate the likely replication of overall model invariance with the assumption that the data used for the estimation represents the larger population. The following arguments are used:

1) `saved_configural`: a saved fitted model at the configural level with no equality constraints. This model should include all other lavaan settings you would like to use, such as estimator or ordered.

2) `data`: The dataframe where the model was estimated.

325   3) `model`: The model syntax for the overall model.

326   4) `group`: The grouping variable column in the dataframe.

327   5) `nboot`: The number of bootstraps to run.

328   6) `invariance_index`: The fit index you would like to use to determine invariance.

329   Please use options and labeling from *lavaan* - see `fitmeasures()` for options.

330   7) `invariance_rule`: The invariance difference score you would like to use as your rule.

331   8) `group.equal`: The equality constraints including in your original multigroup tests.

332   The data included in this function will be sampled, with replacement, at the same

333   size as the current dataset, and the included invariance equality constraints are estimated.

334   Each step will be compared to the previous step using the invariance index and comparison

335   rule entered. The output is a dataframe of the proportion of non-invariant bootstraps from

336   the real data and the same bootstrapped dataset with the group labels randomly assigned.

337   The effect size comparison of proportions, h, for non-invariant comparisons:

$$h_{nmi} = 2 \times (asin\sqrt{p_{data}} - asin\sqrt{p_{random}})$$

338   The alternative, h_{mi}, for effect size of measurement invariance replication would

339   simply be the inverse sign of h_{nmi} and is also included in the table. Two additional columns

340   h_{nmi_p} and h_{nmi_p} represent the h values divided by the upper bound of h (i.e., π), to help

341   with interpretation of the effect size (thus, bounding h to -1 to 1).

342   **Parameter Replication and Effect Sizes: `bootstrap_partial()` Function**

343   After examining the overall model potential replication effect size, the individual

344   parameters within a model can be bootstrapped for partial invariance to with that

345   parameter relaxed (overall partial model statistics) and the difference in group parameter

346   estimates (parameter effect size). This function uses arguments seen in other functions, so

347   they will not be repeated here. The general setup consists of using the model you think

348 could be partially invariant in the `saved_model` argument and the fit index for comparison

349 for the model with less constraints in `invariance_compare`. The `partial_step` argument

350 will be used to determine which operation syntax (i.e. `=~` for loadings) to relax for modeling.

351   The saved output includes several dataframes and plots. The first is the `boot_DF`

352 which the summary of each run in a dataframe for plotting or summarization. This

353 dataframe includes the estimate for each paramter (`term`) separated by group and type

354 (`boot_1`, `boot_2` are the bootstrapped estimates for group 1 and group 2, while the same

355 `random` columns indicate the randomly assigned groups). The fit index used to determine

356 invariance is included for bootstrapped and random estimates, and then the differences

357 between groups and if they were "invariant" or not given the researcher supplied rule.

358   Next, the `boot_summary` includes a summarized form of the bootstrapped results

359 from separated by bootstrapping versus random and invariant/non-invariant. The d_s for

360 between groups Cohen's d is included. Effect sizes are only calculated when the number of

361 bootstrapped estimates is at least 10% of the data - therefore, you would not receive effect

362 sizes with almost no bootstrapped runs. This dataframe should be used to determine which

363 parameter may be different and at what size between groups in a replication of the study.

364 The `boot_effects` table creates a summary similar to the overall model replication table

365 based on the proportion of runs that were considered invariant versus not for each parameter.

366   Plots of the results from dataframes can be found within the `bootstrap_partial()`

367 function. Figure 2 shows the difference between parameters for groups in the bootstrapped

368 and randomly assigned group runs in simulated data. Figure 3 shows the density plot of the

369 estimates for each group organized by bootstrapped and randomly assigned groups and the

370 invariance decision for each bootstrapped run. Last, Figure 4 indicates the d_s value between

371 groups with an indication of the number of data points in each estimate (i.e., dot size).

372 These visualizations should allow a researcher to understand the likelihood of replication for

373 each parameter, as well as the potential size of the differences. Therefore, one could indicate

³⁷⁴ a specific smallest effect size of interest, rather than a invariance cut-off rule of thumb when

³⁷⁵ planning a replication or registered report.

Figure 2

Visualization of the difference score between groups by parameter for invariant and non-invariant bootstrapped and randomly assigned group data.

³⁷⁶ <center>**An Example Analysis**</center>

³⁷⁷   Aiena et al. (2014) examined the RS-14 (Wagnild, 2009) exploring the factor

³⁷⁸ structure of the Resiliency Scale in a clinical sample receiving treatment services and a

³⁷⁹ college student sample. Measurement invariance was calculated for differences separately for

³⁸⁰ these samples for gender and race finding a partially invariant models with a few item

³⁸¹ intercepts or residuals that differed between groups. Aiena et al. (2014) did not compare the

³⁸² clinical to the student sample for measurement invariance, and it is reasonable to expect

³⁸³ potential differences in these two populations. This example will demonstrate the procedure

³⁸⁴ for researchers who wish to use partial invariance steps and how to interpret real, messy data.

³⁸⁵   Table 1 indicates the results after running the one-factor model. There are several

³⁸⁶ guidelines for assessing assessing a degradation in model fit (Cao & Liang, 2022; Cheung &

³⁸⁷ Rensvold, 2002; Counsell et al., 2020; Jin, 2020; Putnick & Bornstein, 2016) but for the

³⁸⁸ purposes of this illustration ΔCFI > .01 will be used. Table 1 indicates that fit was degraded

³⁸⁹ when the constraint on equal item intercepts was added. The code below provides an

³⁹⁰ example of testing each item individually by relaxing the constraints and recalculating the

Figure 3

Visualization of the number of estimates for each group by bootstrapped and randomly assigned group runs by their invariance decision.

Figure 4

Visualization of effect size between groups by parameter for invariant and non-invariant bootstrapped and randomly assigned group data. The size of the dots indicate the number of data points for that estimate.

Table 1

Model Fit for RS-14 Example

Model	AIC	BIC	CFI	TLI	RMSEA	SRMR
Overall	126,722.491	126,888.707	0.934	0.923	0.094	0.036
Group Clinical	52,961.421	53,099.720	0.919	0.904	0.090	0.044
Group Student	69,100.985	69,254.310	0.928	0.915	0.108	0.035
Configural	122,118.406	122,617.055	0.926	0.912	0.102	0.036
Loadings	122,144.532	122,566.010	0.925	0.918	0.098	0.043
Intercepts	122,544.109	122,888.415	0.911	0.910	0.103	0.052
Residuals	126,466.241	126,727.438	0.780	0.793	0.156	0.086

391 CFI. If these Items bring the CFI value back up to ΔCFI $<=$ .01 from the metric model,

392 then the model would be considering partially invariant at the scalar level. It seems unlikely

393 that the residuals will show invariance, if partial scalar invariance can be found, as the drop

394 in fit is quite large.

395   The output indicates that RS6 and RS7 are potential items that could be relaxed to

396 improve model fit and create a partial scalar invariant model (i.e., by picking the largest CFI

397 values). The code below show to check the addition of these items, which are added one at a

398 time. You use the `group.partial` open to "relax" or freely estimate that parameter for each

399 group separately.

400   By examining our estimates, we can see that item seven on the RS-14 is estimated at

401 nearly 5 points for the clinical sample, while the student sample has a lower mean around 4.5

402 points. Generally, students show higher means on the items of the RS-14, but when all

403 loadings and other intercepts are constrained to be equal, and this one item is relaxed, this

404 pattern flips so that clinical groups show higher item intercepts. Given the scale is a 1-7

405 Likert type scale, .5 a point represents a potentially sizable change on the scale. Item seven

⁴⁰⁶ covers perseverance after hardship, and all items can be found in the user manual for the

⁴⁰⁷ scale at www.resiliencecenter.com. The effect size from d_{MACS} suggests a small to medium

⁴⁰⁸ effect, 0.28. In this next code section, we repeat this process for the RS6, as the CFI for our

⁴⁰⁹ model with only RS7 does not achieve the levels of partial invariance for our ΔCFI criterion

⁴¹⁰ (i.e., <= .01 downward change in fit: metric CFI = .925, partial scalar CFI = .914). See

⁴¹¹ Figure 5 for the difference between item intercepts and latent means.

⁴¹²   Again, we see about a half-point difference between our clinical and student samples

⁴¹³ for item 6, which is about drive to achieve. The CFI for this model does meet the

⁴¹⁴ requirements for partial invariance, .917. The effect size is approximately the same at d_{MACS}

⁴¹⁵ = 0.28. See Figure 6 shows the difference between item intercepts and latent means.

Figure 5

RS7 Item Invariance Visualization

Figure 6

RS6 Item Invariance Visualization

Next, we would examine our replication potential for this model. Given our current results, we may not expect our intercepts to replicate. Given the order of `group.equal`, the boot function will select the first non-invariant step in the calculation of the effect size for potential replication. In our output, we do not see a loadings effect size, and this result occurs when *none* of the bootstrapped or random results are non-invariant. Therefore, we would expect the loadings to replicate (and the effect size would be 0 difference between bootstrapped and random, both showing invariance). The intercepts show a large (i.e., close to the max possible value) non-invariant effect, and therefore, we should not expect this model to show invariance in a replication.

Next, we would examine the strength of the effects of replication on each parameter

at the intercept level. By examining Table 2, it is clear that most of the item means are

unlikely to replicate, even though two particular items can be used to create partial

invariance. Figures 7 and 8 display the three plots provided in the `bootstrap_partial()`

function. In general, we should expect $M_D = 0.11$ when items are invariant and $M_D = 0.25$

when items are not invariant. The effect size of non-invariant items ranges from 0.43 to 0.59.

The density plot shown at the bottom of Figure 7 illustrates the likely reasons for the

differences found in the top plots. It appears that many items show a bimodal distribution

within group 1 (Clinical Sample) and when items are invariant, the intercept averages to the

same intercept as group 2 (Student Sample). In non-invariant estimates, the same bimodal

distributions are found, but they are more extreme than the student samples, and therefore,

item show different averages due to the presence of two separate means of data. Further,

some items also appear to show two separate student item averages within the data. This

result suggests that it would be fruitful to understand a potential predictor of these

differences or other confounding variable that separates these samples, creating differences in

item averages.

In summary, if one were planning a replication, the prediction would be that item

intercepts would likely not replicate, with a large effect size (i.e., it is easy to judge h_{nmi_p}

close to the max of one as large). While this study found partial invariance by relaxing

constraints on two individual items, bootstrapped partial invariance indicates that any item

could potentially be problematic with an effect size averaging $d \sim 0.50$ difference in means.

While d_{MACS} values represented a "small" effect based on previous publications, this effect

may be muted by examining both loadings and intercepts. The results here suggest that the

effect is driven by intercepts. The overall average score on items is high: $M_M = 5.04$ (M_{SD}

$= 1.72$). Given the mean standard deviation, a $d \sim 0.50$ represents 0.86 or nearly one whole

point on the scale. A researcher could decide that at least $d = 0.33$ or at least a third of a

standard deviation would be an important change and set that as their smallest effect size of

Table 2

Boot Partial Effects Results for RS-14 Intercepts

Term	Non-Invariant	Random Non-Invariant	h_{nmi}	h_{nmi_p}
RS Intercept	0.989	0.013	2.703	0.860
RS1 Intercept	0.979	0.013	2.622	0.835
RS10 Intercept	0.981	0.013	2.636	0.839
RS11 Intercept	0.988	0.013	2.694	0.857
RS12 Intercept	0.986	0.013	2.676	0.852
RS13 Intercept	0.984	0.013	2.659	0.847
RS14 Intercept	0.986	0.013	2.676	0.852
RS2 Intercept	0.971	0.012	2.580	0.821
RS3 Intercept	0.981	0.012	2.646	0.842
RS4 Intercept	0.989	0.012	2.712	0.863
RS5 Intercept	0.976	0.013	2.602	0.828
RS6 Intercept	0.970	0.013	2.565	0.816
RS7 Intercept	0.961	0.013	2.515	0.801
RS8 Intercept	0.986	0.013	2.676	0.852
RS9 Intercept	0.978	0.013	2.615	0.832

interest for invariance. Further, a newly planned study should investigate what variables may predict when and why samples separate into bimodal representations for item means.

An Example Extension

One benefit of the open science movement on scale development is the publication of datasets or covariance tables with published articles. We can extend our examination of potential replication on other variables that may effect assessment of underlying phenomena. For example, scale translation across languages is not only impacted by the literal conversion
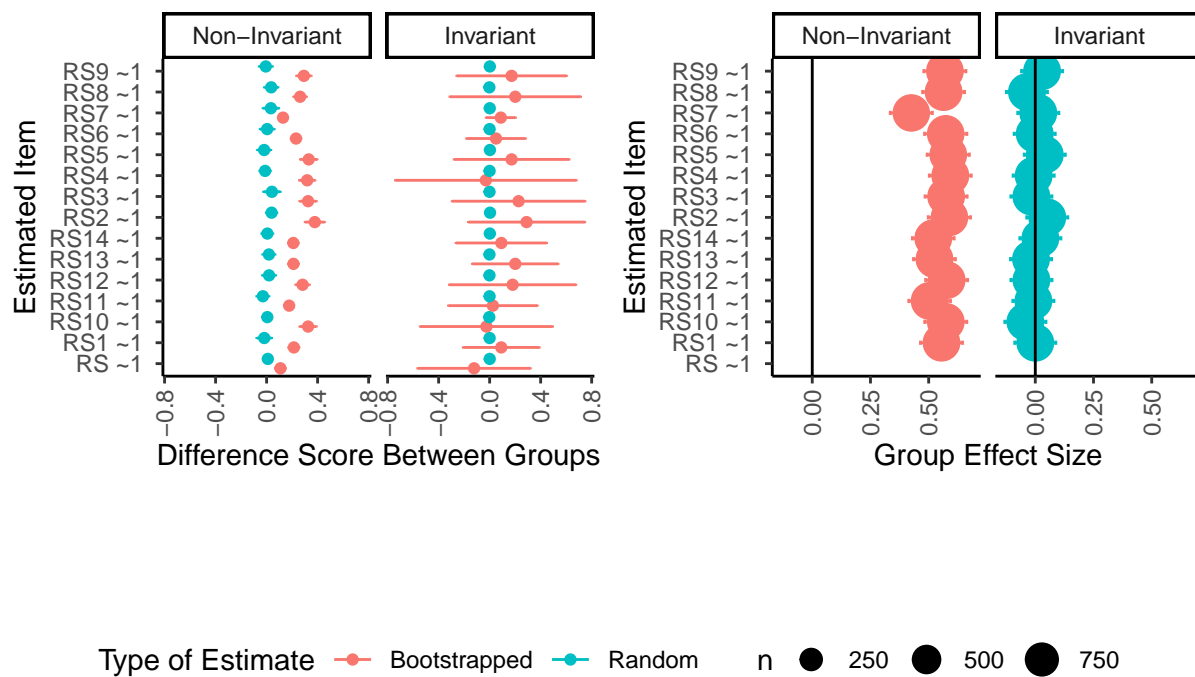
Figure 7

RS-14 scale invariance for item intercepts. The left panel indicates the raw score difference between groups and items, while the right panel indicates the effect size for group differences based on invariance.

459 of concepts from one language to another, but also the cultural and societal norms of the

460 target population (Cha et al., 2007; Chang et al., 1999; Yu et al., 2004). The RS14 was

461 tested in Chinese speaking samples in Chen et al. (2020) across five different large samples

462 and determined that the scale showed good psychometric properties for use within Chinese

463 speaking samples as a one-factor model of resiliency. Given these results, another researcher

464 may assume that the models would be comparable between English speaking (i.e., likely

465 United States) and Chinese speaking samples. With the published data, we can use

466 `visualizemi` to determine if the RS14 is invariant across language/culture, and if the results

467 would likely replicate if tested on a new set of samples, and what, if any, effect size differences
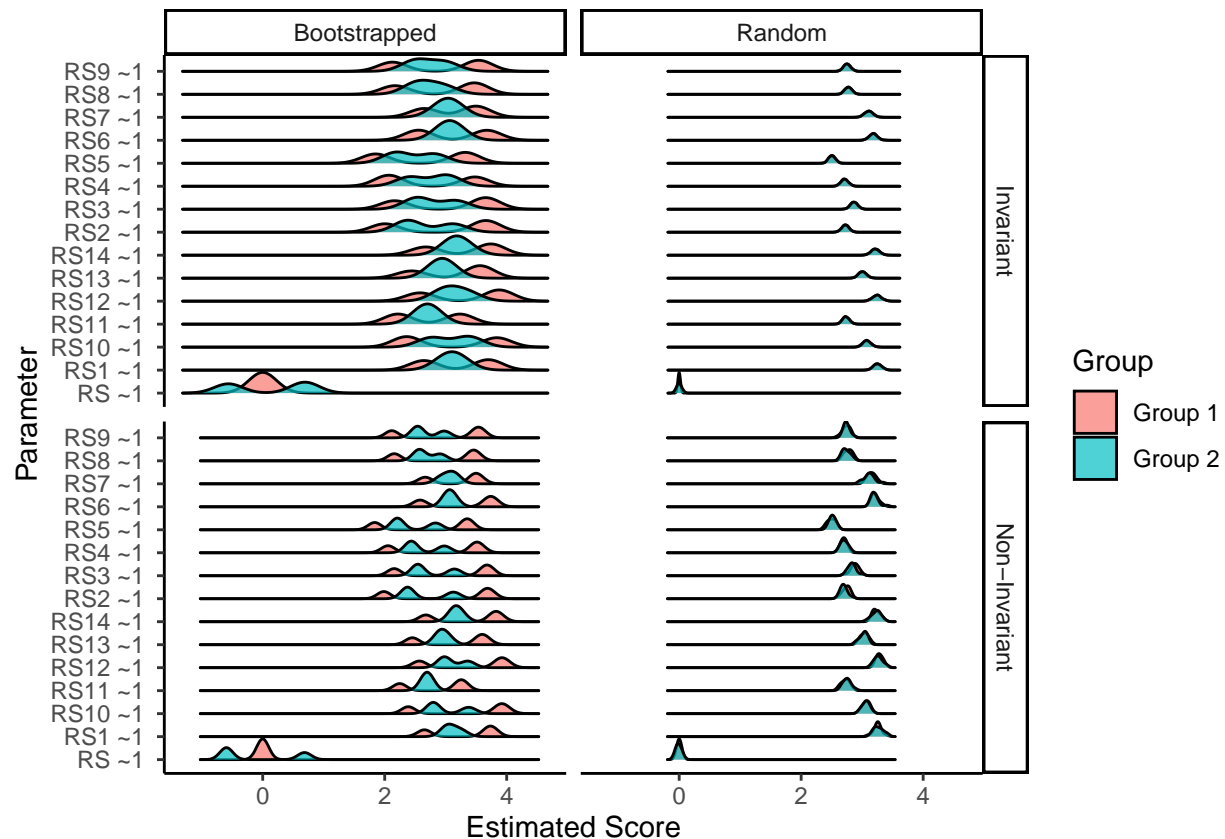
Figure 8

RS-14 scale invariance density plots, illustrating invariant versus non-invariant bootstrapped and random runs for each parameter.

468   are found in parameters. The code used for this analysis is presented in Appendix 7.

469   The Aiena et al. (2014) data used in our previous example was first filtered to only

470   college students, as we have already noted that clinical samples show slightly different

471   intercepts for at least two of the items from student samples. The Chen et al. (2020) data

472   also included college students, which allows us to test a comparable sample that varies on

473   translation and culture. The test of measurement invariance indicated that the factor

474   structure and loadings were invariant across groups. Much like our test of clinical versus

475   student samples, the results indicated that the intercepts were not consistent across groups.

476   Within the English clinical/student sample, partial invariance could be achieved by relaxing

two item intercept constraints (item 6 and 7). To achieve "partial" invariance between the Chinese and English speaking samples, we would need to relax more than half of the items (specifically, eight items: 1, 2, 3, 4, 10, 11, 12, 14), and it would be difficult to suggest partial invariance given this finding. d_{MACS} values range from 0.26 to 0.32 for the eight items which could be interpreted a small to medium given Nye et al. (2019)'s simulation study. Figure 9 portrays the results from the second item on the RS14 (*life accomplishments*).

The results of model bootstrapping indicated that the effect of the loadings was likely to replicate (only invariant results were found), but the intercepts were never found to be invariant compared to a randomized sample ($h_{nmi_p} = 1$). Therefore, we would expect that these differences are persistent, either due to the adaptation or cultural differences in resiliency across samples. The bootstrapped partial invariance demonstrates that each item intercept has a medium to large difference between the two samples as shown in Figure 10, which may explain why full or partial invariance cannot be achieved. This result does not invalidate either version of the scale, but informs researchers of potential baseline differences in item responding for the two samples. Therefore, careful interpretation should be made when comparing these two samples in other instances, as differences in latent means may be the default finding, but with these results one may determine if their results are different from expectations of general scale responding.

Discussion

In this tutorial, we examined how to use multiple tools to examine measurement invariance and its potential replication. Model fit comparisons and statistics can be paired with the proposed effect size measures, and a visualization to examine individual items and the overall latent mean scores. The impact of potential replication was estimated on the overall model and the individual parameters. Using real data, the effect of two non-invariant item intercepts was examined and visualized. This tutorial manuscript has provided a concrete way to plan for pre-registration and/or registered reports. Researchers could

Figure 9

The differences in intercepts for the second item on the RS14 by language sample. The differences between intercepts are shown on the left-hand side with a clear separation between lines. The latent means also show a clear difference between groups where the English group appears to have lower scores overall than the Chinese group.

simulate results based on published or previously collected data to determine the likelihood and size of potential replication. They could plan and pre-register a smallest effect of interest. For example, we may determine that an h_{nmi_p} value above .20 represents an important level of non-invariance for our model overall, while $h_{nmi_p} > .30$ for any individual parameter warrant caution against invariance for groups. Others have begun to discuss the importance of focusing on effects in the scale of the data and their practical importance (Anvari & Lakens, 2021; Cumming, 2012).

Figure 10

Effect size differences for item intercepts on the English versus Chinese samples for the RS14.

⁵¹⁰   From the example, our interpretation may be that the difference between group's

⁵¹¹ latent means is large, as a 0.72 change on a 7 point scale is approximately 10% more

⁵¹² resiliency for students when compared to the clinical sample. Practically, 10% in resiliency

⁵¹³ for an area of the United States (Mississippi) often hit with natural disasters (hurricanes,

⁵¹⁴ tornadoes, floods) and high levels of poverty would be very important. Even the smaller

⁵¹⁵ difference of .5 point on each individual item could translate into increases in resiliency, and

⁵¹⁶ these results may elucidate avenues for further exploration into areas of focus within

⁵¹⁷ resiliency, given the items. The secondary example showed that we can extend these results

⁵¹⁸ to other samples to examine other potentially impactful variables on assessment. The

⁵¹⁹ findings replicate in the sense that the scale shows the same invariant issues with intercepts

⁵²⁰ on a Chinese versus English sample comparison. However, in this analysis, it is clear that the

⁵²¹ differences in items averages are much larger and across all items, rather than a few.

⁵²²   What do the results of a study on measurement invariance with these results tell us

⁵²³ about replication, generalizability, and validity overall? If a researcher decides their effects

⁵²⁴ are large, they should likely caution against suggesting that these scores are directly

⁵²⁵ comparable without weighting or other adjustment. Let's consider a scenario wherein the

⁵²⁶ change metric between models picked (i.e., ΔCFI, ΔRMSEA) indicates a "significant"

⁵²⁷ change in model fit. However, if both the effect size and a visual inspection of the invariance

⁵²⁸ indicates a small difference, we may decide to lessen the practical importance of those results,

⁵²⁹ much like "just significant" p-values with small effect sizes are treated now. The results from

⁵³⁰ our Chinese versus English comparison show us the other scenario: non-invariant results that

⁵³¹ clearly indicate differences with a large effect size on both replication and item average

⁵³² differences. Given that the goal of measurement invariance is to compare *estimates*, we

⁵³³ should expect some differences across samples due to the nature of sampling and estimation.

⁵³⁴ It may be that many of the published models presented represent these effects - small

⁵³⁵ variations between groups due to sampling error or other small crud - but do not represent a

⁵³⁶ fundamental problem with the measurement or generalizability of the results. The

⁵³⁷ `visualizemi` package is one useful tool to help sort out if findings are this small sampling

⁵³⁸ error or differences in samples due to other variables.

References

Aiena, B. J., Baczwaski, B. J., Schulenberg, S. E., & Buchanan, E. M. (2014). Measuring Resilience With the RS–14: A Tale of Two Samples. *Journal of Personality Assessment*, *97*(3), 291–300. https://doi.org/10.1080/00223891.2014.951445

Akaike, H. (1998). *Information theory and an extension of the maximum likelihood principle* (E. Parzen, K. Tanabe, & G. Kitagawa, Eds.; pp. 199–213). Springer New York. http://link.springer.com/10.1007/978-1-4612-1694-0_15

Anvari, F., & Lakens, D. (2021). Using anchor-based methods to determine the smallest effect size of interest. *Journal of Experimental Social Psychology*, *96*, 104159. https://doi.org/10.1016/j.jesp.2021.104159

Barry, A. E., Chaney, B., Piazza-Gardner, A. K., & Chavarria, E. A. (2014). Validity and Reliability Reporting Practices in the Field of Health Education and Behavior: A Review of Seven Journals. *Health Education & Behavior*, *41*(1), 12–18. https://doi.org/10.1177/1090198113483139

Beaujean, A. A. (2014). *Latent variable modeling using r: A step by step guide.* Routledge/Taylor & Francis Group.

Bentler, P. M. (1990). Comparative fit indexes in structural models. *Psychological Bulletin*, *107*(2), 238–246. https://doi.org/10.1037/0033-2909.107.2.238

Bentler, P. M. (1995). *EQS structural equations program manual.*

Boker, S., Neale, M., Maes, H., Wilde, M., Spiegel, M., Brick, T., Spies, J., Estabrook, R., Kenny, S., Bates, T., Mehta, P., & Fox, J. (2011). OpenMx: An Open Source Extended Structural Equation Modeling Framework. *Psychometrika*, *76*(2), 306–317. https://doi.org/10.1007/s11336-010-9200-6

Brown, T. A. (2015). *Confirmatory factor analysis for applied research* (Second edition). The Guilford Press.

Byrne, B. M., Shavelson, R. J., & Muthén, B. (1989). Testing for the equivalence of factor covariance and mean structures: The issue of partial measurement invariance.

Psychological Bulletin, *105*(3), 456–466. https://doi.org/10.1037/0033-2909.105.3.456

Cao, C., & Liang, X. (2022). The impact of model size on the sensitivity of fit measures in measurement invariance testing. *Structural Equation Modeling: A Multidisciplinary Journal*, *29*(5), 744–754. https://doi.org/10.1080/10705511.2022.2056893

Cha, E.-S., Kim, K. H., & Erlen, J. A. (2007). Translation of scales in cross-cultural research: issues and techniques. *Journal of Advanced Nursing*, *58*(4), 386–395. https://doi.org/10.1111/j.1365-2648.2007.04242.x

Chang, A. M., Chau, J. P. C., & Holroyd, E. (1999). Translation of questionnaires and issues of equivalence. *Journal of Advanced Nursing*, *29*(2), 316–322. https://doi.org/10.1046/j.1365-2648.1999.00891.x

Chen, W., Xie, E., Tian, X., & Zhang, G. (2020). Psychometric properties of the Chinese version of the Resilience Scale (RS-14): Preliminary results. *PLOS ONE*, *15*(10), e0241606. https://doi.org/10.1371/journal.pone.0241606

Cheung, G. W., & Rensvold, R. B. (2002). Evaluating Goodness-of-Fit Indexes for Testing Measurement Invariance. *Structural Equation Modeling: A Multidisciplinary Journal*, *9*(2), 233–255. https://doi.org/10.1207/s15328007sem0902_5

Chorpita, B. F., Yim, L., Moffitt, C., Umemoto, L. A., & Francis, S. E. (2000). Assessment of symptoms of DSM-IV anxiety and depression in children: a revised child anxiety and depression scale. *Behaviour Research and Therapy*, *38*(8), 835–855. https://doi.org/10.1016/S0005-7967(99)00130-8

Cohen, J. (2013). *Statistical Power Analysis for the Behavioral Sciences*. Routledge. https://doi.org/10.4324/9780203771587

Cohen, J., Cohen, P., West, S. G., & Aiken, L. (2003). *Applied multiple regression / correlation analysis for the behavioral sciences* (3rd ed.). Lawrence Erlbaum Associates.

Counsell, A., Cribbie, R. A., & Flora, D. B. (2020). Evaluating equivalence testing methods for measurement invariance. *Multivariate Behavioral Research*, *55*(2), 312–328. https://doi.org/10.1080/00273171.2019.1633617

Cumming, G. (2012). *Understanding the new statistics: Effect sizes, confidence intervals, and meta-analysis.* Routledge.

Cumming, G., & Calin-Jageman, R. (2016). *Introduction to the New Statistics* (0th ed.). Routledge. https://doi.org/10.4324/9781315708607

DeVellis, R. F., & Thorpe, C. T. (2022). *Scale development: Theory and applications* (Fifth edition). SAGE Publications, Inc.

Dueber, D. (2023). *Dmacs.* https://github.com/ddueber/dmacs

Epskamp, S. (2022). *semPlot: Path diagrams and visual analysis of various SEM packages' output.* https://CRAN.R-project.org/package=semPlot

Flake, J. K., & Fried, E. I. (2020). Measurement Schmeasurement: Questionable Measurement Practices and How to Avoid Them. *Advances in Methods and Practices in Psychological Science*, *3*(4), 456–465. https://doi.org/10.1177/2515245920952393

Hobson, H. (2019). Registered reports are an ally to early career researchers. *Nature Human Behaviour*, *3*(10), 1010–1010. https://doi.org/10.1038/s41562-019-0701-8

Jin, Y. (2020). A note on the cutoff values of alternative fit indices to evaluate measurement invariance for ESEM models. *International Journal of Behavioral Development*, *44*(2), 166–174. https://doi.org/10.1177/0165025419866911

Jöreskog, K. G. (1971). Simultaneous factor analysis in several populations. *Psychometrika*, *36*(4), 409–426. https://doi.org/10.1007/BF02291366

Jöreskog, K. G., & Sörbom, D. (2001). *LISREL 8: user's reference guide* (2. ed., updated to LISREL 8). SSI Scientific Software Internat.

Lakens, D. (2017). Equivalence Tests. *Social Psychological and Personality Science*, *8*(4), 355–362. https://doi.org/10.1177/1948550617697177

Makel, M. C., & Plucker, J. A. (2014). Facts Are More Important Than Novelty: Replication in the Education Sciences. *Educational Researcher*, *43*(6), 304–316. https://doi.org/10.3102/0013189X14545513

Makel, M. C., Plucker, J. A., & Hegarty, B. (2012). Replications in Psychology Research:

How Often Do They Really Occur? *Perspectives on Psychological Science*, *7*(6), 537–542. https://doi.org/10.1177/1745691612460688

Marsh, H. W., Hau, K.-T., & Wen, Z. (2004). In search of golden rules: Comment on hypothesis-testing approaches to setting cutoff values for fit indexes and dangers in overgeneralizing hu and bentler's (1999) findings. *Structural Equation Modeling: A Multidisciplinary Journal*, *11*(3), 320–341. https://doi.org/10.1207/s15328007sem1103_2

Mayo-Wilson, E., Grant, S., Supplee, L., Kianersi, S., Amin, A., DeHaven, A., & Mellor, D. (2021). Evaluating implementation of the transparency and openness promotion (TOP) guidelines: The TRUST process for rating journal policies, procedures, and practices. *Research Integrity and Peer Review*, *6*(1), 9. https://doi.org/10.1186/s41073-021-00112-8

Meredith, W. (1993). Measurement invariance, factor analysis and factorial invariance. *Psychometrika*, *58*(4), 525–543. https://doi.org/10.1007/BF02294825

Nelson, L. D., Simmons, J., & Simonsohn, U. (2018). Psychology's renaissance. *Annual Review of Psychology*, *69*(1), 511–534. https://doi.org/10.1146/annurev-psych-122216-011836

Nosek, B. A., Ebersole, C. R., DeHaven, A. C., & Mellor, D. T. (2018). The preregistration revolution. *Proceedings of the National Academy of Sciences of the United States of America*, *115*(11), 2600–2606. https://doi.org/10.1073/pnas.1708274114

Nosek, B. A., & Lakens, D. (2014). Registered Reports: A Method to Increase the Credibility of Published Results. *Social Psychology*, *45*(3), 137–141. https://doi.org/10.1027/1864-9335/a000192

Nye, C. D., Bradburn, J., Olenick, J., Bialko, C., & Drasgow, F. (2019). How Big Are My Effects? Examining the Magnitude of Effect Sizes in Studies of Measurement Equivalence. *Organizational Research Methods*, *22*(3), 678–709. https://doi.org/10.1177/1094428118761122

Nye, C. D., & Drasgow, F. (2011). Effect size indices for analyses of measurement equivalence: Understanding the practical importance of differences between groups.

Journal of Applied Psychology, 96(5), 966–980. https://doi.org/10.1037/a0022955

Putnick, D. L., & Bornstein, M. H. (2016). Measurement invariance conventions and reporting: The state of the art and future directions for psychological research. *Developmental Review, 41*, 71–90. https://doi.org/10.1016/j.dr.2016.06.004

Robinson, D., Hayes, A., Couch [aut, S., cre, Software, P., PBC, Patil, I., Chiu, D., Gomez, M., Demeshev, B., Menne, D., Nutter, B., Johnston, L., Bolker, B., Briatte, F., Arnold, J., Gabry, J., Selzer, L., Simpson, G., . . . Reinhart, A. (2023). *Broom: Convert statistical objects into tidy tibbles.* https://CRAN.R-project.org/package=broom

Rosseel, Y. (2012). Lavaan: An r package for structural equation modeling. *Journal of Statistical Software, 48*(1), 1–36. https://doi.org/10.18637/jss.v048.i02

Schwarz, G. (1978). Estimating the dimension of a model. *The Annals of Statistics, 6*(2), 461–464. https://www.jstor.org/stable/2958889

Shadish, W. R., Cook, T. D., & Campbell, D. T. (2001). *Experimental and quasi-experimental designs for generalized causal inference.* Houghton Mifflin.

Sörbom, D. (1978). An alternative to the methodology for analysis of covariance. *Psychometrika, 43*(3), 381–396. https://doi.org/10.1007/BF02293647

Stark, S., Chernyshenko, O. S., & Drasgow, F. (2004). Examining the Effects of Differential Item (Functioning and Differential) Test Functioning on Selection Decisions: When Are Statistically Significant Effects Practically Important? *Journal of Applied Psychology, 89*(3), 497–508. https://doi.org/10.1037/0021-9010.89.3.497

Stark, S., Chernyshenko, O. S., & Drasgow, F. (2006). Detecting differential item functioning with confirmatory factor analysis and item response theory: Toward a unified strategy. *Journal of Applied Psychology, 91*(6), 1292–1306. https://doi.org/10.1037/0021-9010.91.6.1292

Steiger, J. H. (1990). Structural model evaluation and modification: An interval estimation approach. *Multivariate Behavioral Research, 25*(2), 173–180. https://doi.org/10.1207/s15327906mbr2502_4

674 Trent, L. R., Buchanan, E., Ebesutani, C., Ale, C. M., Heiden, L., Hight, T. L., Damon, J.

675   D., & Young, J. (2013). A measurement invariance examination of the revised child

676   anxiety and depression scale in a southern sample: Differential item functioning between

677   african american and caucasian youth. *Assessment*, *20*(2), 175–187.

678   https://doi.org/10.1177/1073191112450907

679 Tucker, L. R., & Lewis, C. (1973). A reliability coefficient for maximum likelihood factor

680   analysis. *Psychometrika*, *38*(1), 1–10. https://doi.org/10.1007/BF02291170

681 Van De Schoot, R., Schmidt, P., De Beuckelaer, A., Lek, K., & Zondervan-Zwijnenburg, M.

682   (2015). Editorial: Measurement invariance. *Frontiers in Psychology*, *6*.

683   https://www.frontiersin.org/articles/10.3389/fpsyg.2015.01064

684 Vazire, S., Schiavone, S. R., & Bottesini, J. G. (2022). Credibility Beyond Replicability:

685   Improving the Four Validities in Psychological Science. *Current Directions in*

686   *Psychological Science*, *31*(2), 162–168. https://doi.org/10.1177/09637214211067779

687 Wagnild, G. (2009). A review of the resilience scale. *Journal of Nursing Measurement*, *17*(2),

688   105–113. https://doi.org/10.1891/1061-3749.17.2.105

689 Weidman, A. C., Steckler, C. M., & Tracy, J. L. (2017). The jingle and jangle of emotion

690   assessment: Imprecise measurement, casual scale usage, and conceptual fuzziness in

691   emotion research. *Emotion*, *17*(2), 267–295. https://doi.org/10.1037/emo0000226

692 Wickham, H. (2016). *ggplot2: Elegant graphics for data analysis*. Springer-Verlag New York.

693   https://ggplot2.tidyverse.org

694 Wilke, C. O. (2020). *Cowplot: Streamlined plot theme and plot annotations for 'ggplot2'*.

695   https://CRAN.R-project.org/package=cowplot

696 Yu, D. S. F., Lee, D. T. F., & Woo, J. (2004). Issues and Challenges of Instrument

697   Translation. *Western Journal of Nursing Research*, *26*(3), 307–320.

698   https://doi.org/10.1177/0193945903260554

699 Zwaan, R. A., Etz, A., Lucas, R. E., & Donnellan, M. B. (2018). Making replication

700   mainstream. *Behavioral and Brain Sciences*, *41*, e120.

Appendix A

Code Examples

702 **Simulating from Models**

703   Here's an example of how to simulate directly from a `lavaan` model:

```r
# first build your model
# this example is separate for each group
model.invariant.g1 <- "
# loadings
lv =~ .8*q1 + .4*q2 + .6*q3 + .3*q4 + .6*q5
# set the residual for invariance on q4
q4 ~~ 1*q4
# set the intercept for invariance on q4
q4 ~ 0*1
# set the intercept to zero for df purposes
q1 ~ 0*1
# allow the latent mean to be estimated
lv ~ 1"
model.invariant.g2 <- "lv =~ .77*q1 + .43*q2 + .58*q3 + .3*q4 + .61*q5
q4 ~~ 1*q4
q4 ~ 0*1
q1 ~ 0*1
lv ~ 1"


# simulate data invariant separately for each group
df.invariant <- bind_rows(
  # lavaan function
  simulateData(
    # model with estimates
```

```r
    model = model.invariant.g1,
    # how many data points
    sample.nobs = 250,
    # mean structure for mgcfa models
    meanstructure = T,
    # model type
    model.type = "cfa",
    # set seed for reproducibility
    seed = 1234) %>%
    # add a group label to the data
    mutate(group = "Group 1"),
  simulateData(
    model = model.invariant.g2,
    sample.nobs = 250,
    meanstructure = T,
    model.type = "cfa",
    seed = 1234) %>%
    mutate(group = "Group 2")
)
```

704 **Simulating from Matrices**

705   Here's an example of how to simulate using `MASS` and covariance or correlation

706 matrices.

```r
library(MASS)

# covariance matrix
university.cov <- lav_matrix_lower2full(
    c(169.00,
      73.710, 182.2500,
      73.229, 88.4250, 171.6100,
      63.375, 72.5625, 127.7250, 156.2500,
      42.120, 67.4325, 122.0265, 123.1875, 182.2500,
      57.226, 63.2610, 117.1926, 154.4250, 138.0240, 201.6400,
      30.875, 32.0625, 60.9805, 62.9375, 76.9500, 79.5910, 90.2500,
      36.075, 38.9610, 61.0722, 58.2750, 65.9340, 70.9290, 81.1965, 123.2100,
      18.096, 21.1410, 26.2131, 39.1500, 44.6310, 46.9452, 48.7635, 56.0106, 75.6900))

# give it names
rownames(university.cov) <-
    colnames(university.cov) <-
    c("class", "social", "learn", "chronic", "physical", "sex",
      "depression", "anxiety", "stress")

# means - you need standard deviation if you only have a correlation matrix
university.means <- c(3.4, 4.3, 3.7, 3.2, 4.5, 1.2, 4.0, 3.5, 4.2)

# use mass function
DF <- mvrnorm(n = 200, mu = university.means, Sigma = university.cov)
```

```r
head(DF)
```

```
##               class       social       learn     chronic    physical         sex
## [1,]    9.332202   6.7442501    8.001600   6.771713 -10.833804  -3.837477
## [2,]    4.873877  -0.1368433   11.690145  14.572699  13.356779  22.699918
## [3,]   13.650244  -5.5841743   -2.295967  -3.914620  -6.552379 -11.711882
## [4,]    7.644520  -1.6201790  -15.075033   4.010138   4.741793  11.704179
## [5,]  -10.491240  14.6367273    2.951522  10.934949   1.153787   3.637487
## [6,]    3.188521  -0.2078648  -11.655781  -8.085560 -12.482893 -18.914448
##          depression    anxiety      stress
## [1,]     -9.2071866  -9.0853468  -10.016408
## [2,]     -5.7266089  -5.1086786   -3.757830
## [3,]    -12.3350189  -9.5855529  -10.855687
## [4,]     -0.8850695  -8.8631134   -1.046865
## [5,]     -3.4707411  -0.1493184   -2.300569
## [6,]     -9.1500394   0.3167367    5.590050
```

MGCFA: `mgcfa()` Function

In this example, we make our example model using *lavaan* syntax. The `lv` latent variable predicts the five measured variables, which are present as columns in our `df.invariant` data set.

lavaan automatically sets the mean (i.e., the intercept) for latent variables to zero. If we wish to visualize the impact of the changes in parameter estimates across groups on the latent means, we need to allow the latent mean estimation with `lv ~ 1`. However, adding this estimation into our model will create a non-identified model. To solve this problem, you can set one of the intercepts of another variable to a value to scale the model. Here we will set the scale of the model by using `q1 ~ 0*1`, thus, scaling the expected means to zero. With simulation, this step is easy to know which variable to pick - we set the intercept on the variable we know did not show differences. In real data, you may wish to run the model steps *without* setting this option, examine the results of a configural or separate models, and then add the option for the values most similar. Additionally, you could complete partial invariance steps to determine which value appears most consistent to fix the estimate.

```
# create lavaan model
model.overall <- "
# overall one-factor model
lv =~ q1 + q2 + q3 + q4 + q5
# set the intercept (mean) of q1 to zero
q1 ~ 0*1
# allow the lv intercept to be freely estimated
lv ~ 1"
# look at the data
head(df.invariant)
```

```
##           q1        q2        q3        q4        q5   group
```

```
## 1 -0.8903542 -0.81707530  0.06137292 -1.3236407 -1.7916418 Group 1
```

```
## 2  1.1054521 -0.03540948 -0.81299606  1.0028340 -0.1909127 Group 1
```

```
## 3  1.4555852  1.54083484  1.59084213 -0.3345967 -0.6865496 Group 1
```

```
## 4 -1.8745187 -1.27880245 -2.53565792 -1.0024193 -1.6253249 Group 1
```

```
## 5 -0.4449517 -0.17782974  1.05507079 -1.2615705  1.7536428 Group 1
```

```
## 6  0.2278813  0.71348845  1.63251893  0.6449847 -1.0055700 Group 1
```

The `mgcfa()` function is designed to flexibly allow you to leverage *lavaan*'s package functions to calculate multiple measurement steps at once. You would include:

1) the model syntax in the `model` argument

2) the dataframe in the `data` argument of our function

3) the name of the grouping variable in quotes for `group`

4) and the equality constraints you would like to impose in order in `group.equal`

5) ... any other *lavaan* arguments you would like to use such as `meanstructure` or `estimator`.

Note: you can also use `sample.cov`, `sample.mean`, `sample.nobs` in this step for estimation of multigroup models, but simulated dataframes are needed for bootstrapping replication estimates.

```r
# run our mgcfa function to run all models
results.invariant <-
  # name of the saved model syntax
  mgcfa(model = model.overall,
        # name of the dataframe
        data = df.invariant,
        # name of the grouping variable
        group = "group",
        # equality constraints to impose in order
```

```
        group.equal = c("loadings", "intercepts", "residuals"),
        # other options to send to lavaan cfa function
        meanstructure = T)


# what is saved for you
names(results.invariant)
```

[1] "model_coef"   "model_fit"   "model_overall"
[4] "group_models"   "model_configural" "invariance_models"

1) `model_coef`: The parameter estimates for each model with the model step included in a *model* column. This set of coefficients can be used for other functions. This dataframe is created with *broom*'s `tidy()` function if you wish to recreate this table without running the `mgcfa()` function (Robinson et al., 2023).

```
results.invariant$model_coef[1:10 , ]
```

```
## # A tibble: 10 x 12
```

##	term	op	estimate	std.error	statistic	p.value	std.lv	std.all	model
##	\<chr\>	\<chr\>	\<dbl\>	\<dbl\>	\<dbl\>	\<dbl\>	\<dbl\>	\<dbl\>	\<chr\>
## 1	"lv =~ q1"	=~	1	0	NA	NA	0.803	0.616	Over~
## 2	"lv =~ q2"	=~	0.655	0.0880	7.44	9.77e-14	0.526	0.493	Over~
## 3	"lv =~ q3"	=~	0.640	0.0895	7.15	8.83e-13	0.514	0.463	Over~
## 4	"lv =~ q4"	=~	0.277	0.0749	3.69	2.24e- 4	0.222	0.209	Over~
## 5	"lv =~ q5"	=~	0.955	0.117	8.13	4.44e-16	0.766	0.656	Over~
## 6	"q1 ~1 "	~1	0	0	NA	NA	0	0	Over~
## 7	"lv ~1 "	~1	-0.0305	0.0582	-0.524	6.00e- 1	-0.0380	-0.0380	Over~
## 8	"q1 ~~ q1"	~~	1.05	0.0995	10.6	0	1.05	0.620	Over~
## 9	"q2 ~~ q2"	~~	0.860	0.0653	13.2	0	0.860	0.757	Over~

10 "q3 ~~ q3" ~~   0.966   0.0711   13.6   0   0.966   0.785   Over~

i 3 more variables: block <int>, group <int>, label <chr>

2) `model_fit`: The model fit indices from `fitmeasures()` to review for overall model fit
and invariance judgments. The name of the model is included in a *model* column.

```
head(results.invariant$model_fit)
```

A tibble: 6 x 18

	agfi	AIC	BIC	cfi	chisq	npar	rmsea	rmsea.conf.high	srmr	tli
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	0.998	7516.	7580.	1	0.650	15	0	0	0.00616	1.04
## 2	0.948	3766.	3819.	0.976	7.79	15	0.0473	0.108	0.0312	0.953
## 3	0.974	3768.	3820.	1	4.48	15	0	0.0831	0.0210	1.01
## 4	0.961	7533.	7660.	0.991	12.3	30	0.0301	0.0785	0.0261	0.982
## 5	0.965	7528.	7638.	0.994	15.4	26	0.0200	0.0660	0.0330	0.992
## 6	0.969	7522.	7615.	1	17.3	22	0	0.0542	0.0352	1.00

i 8 more variables: converged <lgl>, estimator <chr>, ngroups <int>,
missing_method <chr>, nobs <int>, norig <int>, nexcluded <int>, model <chr>

3) `model_overall`: A saved *lavaan* fitted model of all groups together without any
equality constraints or grouping variables. These objects can be used with any function
that normally takes a saved model: `parameterEstimates()`,
`modificationIndices()`, `semPlot::semPaths()`, and so on (Epskamp, 2022).

```
class(results.invariant$model_overall)
```

[1] "lavaan"
attr(,"package")
[1] "lavaan"

794   4) `group_models`: A list of saved fitted models for each group separately.

```r
names(results.invariant$group_models)
```

795 `## [1] "model.Group 1" "model.Group 2"`

796   5) `model_configural`: A saved fitted model for the configural model that nests together

797   each group into one model with no other constraints.

```r
class(results.invariant$model_configural)
```

798 `## [1] "lavaan"`

799 `## attr(,"package")`

800 `## [1] "lavaan"`

801   6) `invariance_models`: A list of saved fitted models that consecutively adds

802   `group.equal` constraints.

```r
names(results.invariant$invariance_models)
```

803 `## [1] "model.loadings"   "model.intercepts" "model.residuals"`

Partial Invariance: `partial_mi()` Function

The `partial_mi()` function aids in the calculation of partial invariance for a specific step of the MGCFA process. The function includes the following arguments:

1) `saved_model`: The saved *lavaan* model with the equality constraints at the level of measurement invariance you would like to examine for partial invariance.

2) `data`: The dataframe where the model was estimated.

3) `model`: The model syntax for the overall model.

4) `group`: The grouping variable column in the dataframe.

5) `group.equal`: The equality constraints including in your original multigroup tests.

6) `partial_step`: The level of partial invariance you wish to test.

```r
partial.invariant <-

  partial_mi(

    # saved model output with constraints

    saved_model = results.invariant$invariance_models$model.residuals,

    # dataframe from model

    data = df.invariant,

    # model syntax

    model = model.overall,

    # group column name

    group = "group",

    # group equality constraints from your mgcfa

    group.equal = c("loadings", "intercepts", "residuals"),

    # which step you want to examine for partial invariance

    partial_step = "residuals"

    )


names(partial.invariant)
```

```
## [1] "models"    "fit_table"
```

In this function, each parameter with the appropriate *lavaan* syntax is relaxed individually (i.e., `~1` for intercepts, `~~` for residuals, etc.). The fitted models are saved in the `models` output, and the `fit_table` output includes all fit indices for each model to investigate potential areas of partial invariance based on the researcher's desired criterion.

```
names(partial.invariant$models)
```

```
## [1] "q1 ~~ q1" "q2 ~~ q2" "q3 ~~ q3" "q4 ~~ q4" "q5 ~~ q5" "lv ~~ lv"
```

```
head(partial.invariant$fit_table %>%
       dplyr::select(free.parameter, cfi, rmsea))
```

```
## # A tibble: 6 x 3
##    free.parameter cfi          rmsea
##    <chr>          <lvn.vctr>   <lvn.vctr>
## 1 q1 ~~ q1        0.9902679    0.02108648
## 2 q2 ~~ q2        0.9868905    0.02447336
## 3 q3 ~~ q3        0.9958241    0.01381266
## 4 q4 ~~ q4        1.0000000    0.00000000
## 5 q5 ~~ q5        0.9868088    0.02454944
## 6 lv ~~ lv        0.9906154    0.02025143
```

Note: the `partial_step` function is used to determine which types of `op` or operators to freely estimate between groups. If one chooses residuals, you will also freely estimate the residual for the latent variable or any other residuals found in the model. These items may be ignored if they were not meant to be included.

Visualization of Invariance: `plot_mi()` Functions

Once we know which items are non-invariant, the `model_coef` output from the `mgcfa()` can be used directly in `plot_mi()`. The plot outputs will be described below. First, here are the arguments for the function:

1) `data_coef`: A tidy dataframe of the parameter estimates from the models. This function assumes you have used `broom::tidy()` on the saved model from *lavaan* and added a column called "model" with the name of the model step (Robinson et al., 2023). This function will only run for models that have used the grouping function (i.e., configural, metric, scalar, and strict or other combinations/steps you wish to examine).

2) `model_step`: Which model do you want to plot? You should match this name to the one you want to extract from your model column in the `data_coef`.

3) `item_name`: Which observed variable from your model syntax do you want to plot? Please list this variable name exactly how it appears in the model.

4) `x_limits`: What do you want the x-axis limits to be for your invariance plot? The default option is to assume the latent variable is standardized, and therefore, -1 to 1 is recommended. Use only two numbers, a lower and upper limit. This value also constrains the latent mean diagram to help zoom in on group differences because the scale of latent means is usually centered over zero. You can use this parameter to zoom out to a more traditional histogram using `c(-2, 2)`.

5) `y_limits`: What do you want the y-axis limits to be for your invariance plot? Given that the latent variable is used to predict the observed values in the data, you could use the minimum and maximum values found in the data. If that range is large, consider reducing this value to be able to visualize the results (i.e., otherwise it may be too zoomed out to judge group differences). Use only two numbers, a lower and upper limit.

6) `conf.level`: What confidence limit do you want to plot? Use $1 - \alpha$.

859 7) `model_results`: In this argument, include the saved *lavaan* output for the model

860   listed in the `model_step` argument.

861 8) `lv_name`: Include the name of the latent variable, exactly how it is listed in your

862   *lavaan* syntax. You should plot the latent variable that the `item_name` is linked to. If

863   you have items that load onto multiple latent variables, you will need to make multiple

864   plots.

865 9) `plot_groups`: If you include more than two groups in a multigroup model, the

866   automatic assumption is that you want the first two groups for this visualization. If

867   not, include the names of the groups here to plot.

```r
invariant.plot <-
  plot_mi(
    # output from model_coef
    data_coef = results.invariant$model_coef,
    # which model do you want to plot
    model_step = "Configural",
    # name of observed item
    item_name = "q4",
    # latent variable limits to graph
    x_limits = c(-1,1),
    # Y min and max in data
    y_limits = c(min(df.invariant$q4), max(df.invariant$q4)),
    # what ci do you want
    conf.level = .95,
    # what model results do you want
    model_results = results.invariant$model_configural,
    # which latent variable do you want
    lv_name = "lv"
)
```

```
names(invariant.plot)
```

[1] "complete" "intercept" "mean"   "variance"

The outputs from this function are several *ggplot2* objects that can be edited or saved directly using *ggplot2* functionality (Wickham, 2016).

1) `complete`: The output from this model can be found in Figure 1. On the left-hand side, the item invariance is plotted, and on the right-hand side, the latent mean distributions for the two groups are plotted. In the item invariance sub-plot, the visualization includes all three components traditionally seen in MGCFA testing steps: loadings, intercepts, and residuals. Each visualization element was designed to match the traditional visualization for that type of output. All parameter estimates are plotted on the unstandardized estimates and their confidence interval based on the standard error of the estimate. All plots are made with *ggplot2* and *cowplot* (Wilke, 2020).

2) `intercept`: Only the left-hand side of the complete plot designed to represent intercepts and factor loadings. Factor loadings represent the slope of the regression equation for the latent variable predicting the scores on the observed variable ($\hat{Y} \sim b_0 + b_1 X + \epsilon$). The y-axis indicates the observed variable scores, and here, the plot includes the entire range of the scale of the data for item four. The coefficient (b_1) for group 1 was 0.40, while the coefficient for group 2 was 0.21. The ribbon bands around the plotted slopes indicate the confidence interval for that estimate. In this plot, while the coefficients for each group are not literally equal, the overlapping and parallel slope bands indicate they are not different practically.

The item intercepts (b_0) are plotted on the middle line where they would cross the y-axis at a latent variable score of zero. These are represented by a dot with a set of confidence error bars around the point. The intercept for group 1 was 0.07, while the

Figure A1

Invariant Model Visualization

891 coefficient for group 2 was 0.03. In this invariant depiction, the overlap in the intercepts is

892 clear, indicating they are not different. You can use `y_limits` to zoom in on the graph if

893 these are too small to be distinguishable.

894 3) `mean`: The right-hand side of the complete plot graphing the latent variable means and

895   density from the data. The latent variable is shown on the x-axis using standardized

896   values (i.e., *z*-scores) where -1 indicates one standard deviation below the mean for the

897   latent variable, 0 indicates the mean for the latent variable and so on. The lines

898   indicate the means of the latent variables from the simulated dataset. Group labels are

899   represented in the figure caption on the bottom. Group 1 is usually the group that is

900   alphabetically first in the data set or whichever group is the first that appears when

⁹⁰¹ using the `levels()` command.

⁹⁰² 4) `variance`: A split geom violin plot indicating the variance distribution of the plotted

⁹⁰³ item. Residuals are trickier to plot, as they are the left over error when predicting the

⁹⁰⁴ observed variables ϵ. It is tempting to plot this value as the confidence band around

⁹⁰⁵ the slope, however, that defeats the purpose of understanding that the slopes are

⁹⁰⁶ estimated separately from the residuals, and both have an associated variability around

⁹⁰⁷ their parameter estimate. Therefore, residuals are represented in the inset picture at

⁹⁰⁸ the bottom right of the item invariance plot. The black bars represent the estimated

⁹⁰⁹ residual for each group (group 1: 0.91, group 2: 1.22). The distributions are plotted to

⁹¹⁰ represent the normal spread of values using the standard error of the residuals. The

⁹¹¹ violin plot allows for direct comparison of those residuals and their potential

⁹¹² distributions. Note that the placement has nothing to do with the x or y-axis and is

⁹¹³ designed to always show in the same location, regardless of size/value. The plots are

⁹¹⁴ included separately so they can be arranged in a different fashion if desired.

Model Replication and Effect Sizes: `bootstrap_model()` Function

The `bootstrap_model` function in *visualizemi* was designed to estimate the likely replication of overall model invariance with the assumption that the data used for the estimation represents the larger population. The following arguments are used:

1) `saved_configural`: a saved fitted model at the configural level with no equality constraints. This model should include all other lavaan settings you would like to use, such as estimator or ordered.

2) `data`: The dataframe where the model was estimated.

3) `model`: The model syntax for the overall model.

4) `group`: The grouping variable column in the dataframe.

5) `nboot`: The number of bootstraps to run.

6) `invariance_index`: The fit index you would like to use to determine invariance. Please use options and labeling from *lavaan* - see `fitmeasures()` for options.

7) `invariance_rule`: The invariance difference score you would like to use as your rule.

8) `group.equal`: The equality constraints including in your original multigroup tests.

```r
boot.model.invariant <-
  bootstrap_model(
    # saved configural model
    saved_configural = results.invariant$model_configural,
    # dataframe
    data = df.invariant,
    # model syntax
    model = model.overall,
    # group variable column in dataframe
    group = "group",
    # number of bootstraps
    nboot = 1000,
```

```
    # which fit index you would like to use

    invariance_index = "cfi",

    # what is your criterion for that fit index

    invariance_rule = .01,

    # what equality constraints are you testing

    group.equal = c("loadings", "intercepts", "resduals")

)
```

930   The data included in this function will be sampled, with replacement, at the same

931 size as the current dataset, and the included invariance equality constraints are estimated.

932 Each step will be compared to the previous step using the invariance index and comparison

933 rule entered. The output is a dataframe of the proportion of non-invariant bootstraps from

934 the real data and the same bootstrapped dataset with the group labels randomly assigned.

935 The effect size comparison of proportions, h, for non-invariant comparisons:

$$h_{nmi} = 2 \times (asin\sqrt{p_{data}} - asin\sqrt{p_{random}})$$

936   The alternative, h_{mi}, for effect size of measurement invariance replication would

937 simply be the inverse sign of h_{nmi} and is also included in the table. Two additional columns

938 h_{nmi_p} and h_{nmi_p} represent the h values divided by the upper bound of h (i.e., π), to help

939 with interpretation of the effect size (thus, bounding h to -1 to 1).

Parameter Replication and Effect Sizes: `bootstrap_partial()` Function

After examining the overall model potential replication effect size, the individual parameters within a model can be bootstrapped for partial invariance to with that parameter relaxed (overall partial model statistics) and the difference in group parameter estimates (parameter effect size). This function uses arguments seen in other functions, so they will not be repeated here. The general setup consists of using the model you think could be partially invariant in the `saved_model` argument and the fit index for comparison for the model with less constraints in `invariance_compare`. This example examines the loadings in the invariant model, so `saved_model` uses the `mgcfa` output for equality constraints present on the loadings and compares that model to the configural model with no equality constraints on the loadings. The `partial_step` argument will be used to determine which operation syntax (i.e. `=~` for loadings) to relax for modeling.

```r
boot.partial.invariant <-
    bootstrap_partial(
        # saved model you want to examine the partial loadings for
        saved_model = results.invariant$invariance_models$model.loadings,
        # the dataset
        data = df.invariant,
        # the model
        model = model.overall,
        # the group variable in the dataset
        group = "group",
        # number of bootstraps
        nboot = 1000,
        # which fit index you would like to use to determine partial invariance
        invariance_index = "cfi",
        # what is the invariance rule
        invariance_rule = .01,
```

```r
        # what are we comparing the saved model fit index to
        invariance_compare = fitmeasures(results.invariant$model_configural, "cfi"),
        # what step are we using for invariance
        partial_step = "loadings",
        # what equality constraints should be imposed
        group.equal = c("loadings")
    )
```

```r
names(boot.partial.invariant)
```

```
## [1] "invariance_plot"      "effect_invariance_plot" "density_plot"
## [4] "boot_DF"              "boot_summary"          "boot_effects"
```

The saved output includes several dataframes and plots. The first is the `boot_DF` which the summary of each run in a dataframe for plotting or summarization. This dataframe includes the estimate for each paramter (`term`) separated by group and type (`boot_1`, `boot_2` are the bootstrapped estimates for group 1 and group 2, while the same `random` columns indicate the randomly assigned groups). The fit index used to determine invariance is included for bootstrapped and random estimates, and then the differences between groups and if they were "invariant" or not given the researcher supplied rule.

```r
head(boot.partial.invariant$boot_DF)
```

```
##          term     boot_1      boot_2 random_1  random_2  boot_fit random_fit
## 1 lv =~ q1 0.4548783 0.49928877 0.4627486 0.4651391 0.9296990  1.0000000
## 2 lv =~ q2 0.3599017 0.56241016 0.4100874 0.4980215 0.9441125  1.0000000
## 3 lv =~ q3 0.4254283 0.33640233 0.4274329 0.3422124 0.9377130  1.0000000
## 4 lv =~ q4 0.3930716 0.03320619 0.1380833 0.2628802 0.9750274  1.0000000
## 5 lv =~ q5 0.7306414 0.73512673 0.7093891 0.7532471 0.9266587  1.0000000
## 6 lv =~ q1 0.5537083 0.57086815 0.5732166 0.5475714 0.8958929  0.9814658
```

```
##    boot_difference random_difference boot_index_difference
## 1     -0.044410454      -0.002390463                 FALSE
## 2     -0.202508484      -0.087934027                 FALSE
## 3      0.089025927       0.085220565                 FALSE
## 4      0.359865463      -0.124796846                 FALSE
## 5     -0.004485377      -0.043857947                 FALSE
## 6     -0.017159815       0.025645271                 FALSE
##   random_index_difference
## 1                    TRUE
## 2                    TRUE
## 3                    TRUE
## 4                    TRUE
## 5                    TRUE
## 6                    TRUE
```

Next, the `boot_summary` includes a summarized form of the bootstrapped results from separated by bootstrapping versus random and invariant/non-invariant. The d_s for between groups Cohen's d is shown below, and the non-central confidence interval is included. Effect sizes are only calculated when the number of bootstrapped estimates is at least 10% of the data - therefore, you would not receive effect sizes with almost no bootstrapped runs. This dataframe should be used to determine which parameter may be different and at what size between groups in a replication of the study.

```
boot.partial.invariant$boot_summary %>%
  dplyr::select(term, d_boot, d_random)
```

```
##         term       d_boot      d_random
## 1   lv =~ q1 -0.029853316   0.058271662
## 2   lv =~ q1  0.033742666   0.011640524
```

```
## 3  lv =~ q2 -0.032613505  0.093288563

## 4  lv =~ q2  0.146200211  0.030925365

## 5  lv =~ q3 -0.046329761  0.112823117

## 6  lv =~ q3 -0.148330246  0.074265298

## 7  lv =~ q4  0.007851687 -0.066844445

## 8  lv =~ q4 -0.015702705  0.038884725

## 9  lv =~ q5 -0.001285809 -0.169307690

## 10 lv =~ q5  0.122405579 -0.008526218
```

The `boot_effects` table creates a summary similar to the overall model replication table based on the proportion of runs that were considered invariant versus not for each parameter. Note that the effects match the overall results, such that simulated invariant data appears to still show the likelihood that loadings may not replicate in a similar dataset.

```
boot.partial.invariant$boot_effects
```

```
##         term non_invariant random_non_invariant    h_nmi     h_mi   h_nmi_p
## 1 lv =~ q1         0.853                0.236 1.340078 -1.340078 0.4265601
## 2 lv =~ q2         0.858                0.237 1.351946 -1.351946 0.4303378
## 3 lv =~ q3         0.851                0.230 1.348639 -1.348639 0.4292851
## 4 lv =~ q4         0.840                0.229 1.320578 -1.320578 0.4203530
## 5 lv =~ q5         0.819                0.237 1.245789 -1.245789 0.3965468
##       h_mi_p
## 1 -0.4265601
## 2 -0.4303378
## 3 -0.4292851
## 4 -0.4203530
## 5 -0.3965468
```

1016   Plots of the results from dataframes can be found within the `bootstrap_partial()`

1017 function. Figure 2 shows the difference between parameters for groups in the bootstrapped

1018 and randomly assigned group runs. Figure 3 shows the density plot of the estimates for each

1019 group organized by bootstrapped and randomly assigned groups and the invariance decision

1020 for each bootstrapped run. Last, Figure 4 indicates the d_s value between groups with an

1021 indication of the number of data points in each estimate (i.e., dot size). These visualizations

1022 should allow a researcher to understand the likelihood of replication for each parameter, as

1023 well as the potential size of the differences. Therefore, one could indicate a specific smallest

1024 effect size of interest, rather than a invariance cut-off rule of thumb when planning a

1025 replication or registered report.
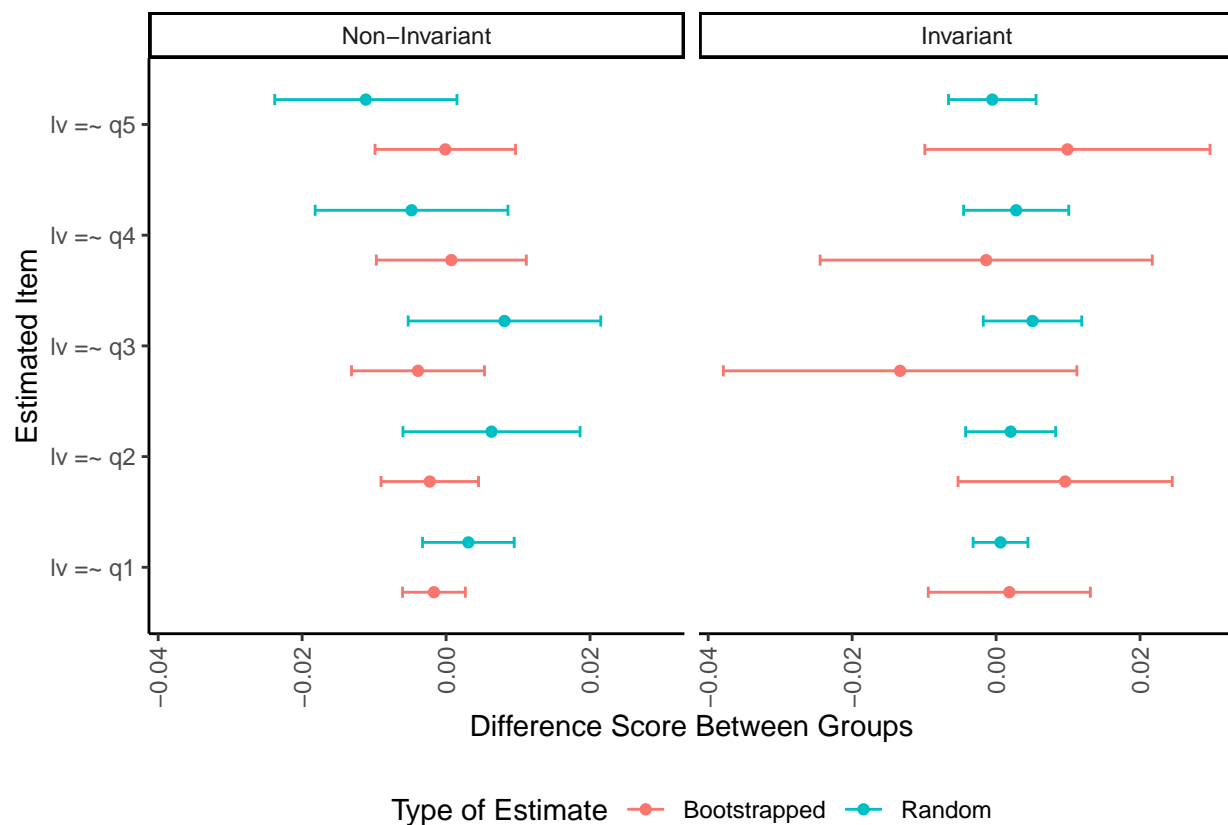
Figure A2

Visualization of the difference score between groups by parameter for invariant and non-invariant bootstrapped and randomly assigned group data.

Figure A3

Visualization of the number of estimates for each group by bootstrapped and randomly assigned group runs by their invariance decision.

Figure A4

Visualization of effect size between groups by parameter for invariant and non-invariant bootstrapped and randomly assigned group data. The size of the dots indicate the number of data points for that estimate.

Appendix B

Simulation Study

Design and Analysis

Data was simulated using the `simulateData` function in the *R* package *lavaan* (Rosseel, 2012) assuming multivariate normality using a μ of 0 and σ of 1 for the data. This function allows you to write *lavaan* syntax for your model with estimated values to generate data for observed variables (see supplemental for examples). The data included two groups of individuals ("Group 1", "Group 2") for a multigroup confirmatory factor analysis (n_{group} = 250, N = 500). The latent variables were assumed to be continuous normal (the package functions do not require this assumption). The model consisted of five observed items predicted by one latent variable (`lv =~ q1 + q2 + q3 + q4 + q5`); however, the demonstration in this manuscript extends to multiple latent variables and other combinations of observed variables. Each item was assumed to be related to the latent variable with loadings approximately equal to .40 to .80, except when cases of non-invariance on the loadings was simulated.

The Brown (2015) steps of testing measurement invariance are demonstrated in this manuscript for illustration purposes, but in line with Stark et al. (2006) suggestions, the visualizations show the impact of loadings and intercepts together. A convenience function `mgcfa` is used for these steps or other measurement invariance test orders and combinations. Fit indices for the steps for multigroup models are presented in the appendix for comparison of cutoff rules of thumb (Cheung & Rensvold, 2002) to effect sizes and visualizations presented in this manuscript. Fit indices include Akaike Information Criterion (AIC, Akaike, 1998), Bayesian Information Criterion (BIC, Schwarz, 1978), Comparative Fit Index (CFI, Bentler, 1990), Tucker Lewis Index (TLI, Tucker & Lewis, 1973), root mean squared error of approximation RMSEA (Steiger, 1990), and standardized root mean square residual (SRMR, Bentler, 1995).

1050 The data was then simulated to represent invariance across all model steps, small,

1051 medium, and large invariance using d_{MACS} estimated sizes from Nye et al. (2019). While

1052 d_{MACS} is used primarily for an effect size of the (non)-invariance for intercepts and loadings

1053 together, a similar approach was taken for the estimation of small, medium, and large effects

1054 on the residuals. The effect size is presented for all models, calculated from the *dmacs*

1055 package (Dueber, 2023; Nye & Drasgow, 2011). Only one item in each model was

1056 manipulated from the invariant model to create the non-invariant models. Given the data

1057 was simulated with a *z*-score scaling, the loading values were simulated at .30 points apart

1058 (given d_{MACS} suggestions of .2, .4, .7), the intercepts at .25 points apart, and the residuals at

1059 .25 points apart. To plan a simulation for your own study, these values can be used to

1060 simulate small, medium, and large non-invariance effects by first converting data into *z*-score.

Visualize Parameter Differences

1062 The d_{MACS} value for item 4 in the invariant model was 0.16, representing a nil or

1063 unimportant difference in this manuscript. It is important to note that while Nye et al.

1064 (2019) suggests specific sizes for small, medium, and large, each researcher should determine

1065 for themselves what effects represent. Figure B1 displays the results from the small (d_{MACS}

1066 = 0.27) difference in loadings, while Figure B2 displays the results from the medium (d_{MACS}

1067 = 0.61) difference in loadings, and Figure B3 shows the large ($d_{MACS} = 0.66$) differences.

1068 When investigating the slope values, we can clearly see the change in the loading for the

1069 second group (the only manipulated variable, although random data set generation may also

1070 change intercepts and residuals slightly). At the medium effect size, we see that the

1071 confidence bands do not overlap (at the edges), and at the large effect size, we can see a clear

1072 separation of two lines. Note that the intercepts in this model are estimated as equal so the

1073 loading representation will not literally separate, but the steepness of the lines is the

1074 indicator of the difference between the slopes. You can imagine these lines are interpreted

1075 like a simple slopes analysis for interactions in regression (Cohen et al., 2003). When simple

1076 slopes for interactions are plotted, if they are parallel, there is no interaction, and if they

1077 cross, then there is an interaction. Here, we can use this same logic. If they are parallel,

1078 there is likely invariance (they are the same), and the further from parallel they become, the

1079 larger the effect size for the differences between group loadings.

1080   The latent means in Figure B3 do appear to show differences, albeit visually small.

1081 The latent means diagram shows the impact of any group differences that aren't constrained,

1082 and this image shows the configural model (as the metric model would force them to be

1083 equal). In the simulated model, the *only* manipulated parameter is question 4's loading. In

1084 real models, the differences may be larger due to other variation found in the parameter

1085 estimates. Therefore, once you discover items you believe would make a model "partially"

1086 invariant, you may wish to estimate that model and graph the item again using the partially

1087 invariant model to see only the effect of the non-invariant items. Additionally, consider that

1088 we set the scaling of the model to 0. The estimate for the lv mean in the large loading model

1089 was group 1: 0.00, and group 2: -0.06, which results in 0.06 difference in group means. The

1090 practical implications of this difference will depend on the research and interpretations of the

1091 researcher.

1092   For intercepts, the small (Figure B4), medium (Figure B5), and large (Figure B6)

1093 depictions represent d_{MACS} values of 0.26, 0.47, and 0.70, respectively. Intercept differences

1094 can be clearly seen represented by the spacing out of the intercept locations (and thus, the

1095 overall line as well). While the changes in intercept do not appear to change the latent

1096 means, the caveat to this simulation is that only item four was manipulated. An example is

1097 provided below that demonstrates large changes in latent means.

1098   Last, the effect of the residuals is plotted in small (Figure B7), medium (Figure B8),

1099 and large (Figure B9) formats. While d_{MACS} values are not technically available for the

1100 residuals, our models showed 0.19, 0.10, and 0.16, respectively. These differences in values

1101 are variable due to the random generation of data sets for each measurement invariance

1102 manipulation. At first glance, the differences in the small chart may seem large, because the

Figure B1

Small Loadings Model Visualization

black lines are not touching, but notice that the distributions overlap, indicating a likely small difference. The medium and large differences better illustrate differences in residuals across groups. Further, the impact of the residuals on the shape of the latent mean distribution can also been seen (and unintentionally, in the first figures as well due to random variation). The impact is due to the standard error of the residuals, as smaller standard errors represent lepokurtic distributions (taller), and larger standard errors represent platykurtic distributions (flatter). The effect size difference of the residuals does not appear to change the effects in the latent means.

Figure B2

Medium Loadings Model Visualization

Model Replication and Effect Size

Figure B10 portrays the h_{nmi_p} values by simulated non-invariance, strength of non-invariance, and type of equality constraint. This image represents 100 simulations of data by 1000 bootstrapped runs (averaged) to explore the expected pattern of results. The bars are arranged to show what a researcher might inspect when thinking about replication possibilities and their effect sizes (i.e., only three bars for each equality constraint would be calculated).

In the data that was simulated to be invariant between groups, effect sizes are still non-zero (loadings $h_{nmi_p} = 0.28$, intercepts $= h_{nmi_p} = 0.06$, $h_{nmi_p} = 0.00$). This result mirrors the effects found in the literature - that often, many models fail to show invariance,

Figure B3

Large Loadings Model Visualization

¹¹²¹ and potentially not because measurement is poor but because of natural random variation in

¹¹²² parameter estimates. This result also indicates the need to be able to identify if specific

¹¹²³ parameters are driving the differences, which is shown in the next section.

¹¹²⁴ Next, Figure B10 demonstrates the patterns one might find for small, medium, and

¹¹²⁵ large effects at each type of invariance when data is simulated with *one* difference. For

¹¹²⁶ loadings, the pattern shows a larger effect for loadings with zero or negative effect sizes for

¹¹²⁷ other effect sizes. The intercept simulations show non-zero effect sizes in the loadings and

¹¹²⁸ intercepts, likely for the same reasons d_{MACS} is interpreted as a combined effect size. When

¹¹²⁹ intercepts are changed, loadings may naturally shift with those means. Last, the residual

¹¹³⁰ results present an unexpected pattern, wherein the effect is primarily seen in the loadings,

Figure B4

Small Intercepts Model Visualization

rather than the residuals step. However, when distributions of error variance are different,

one may expect that those effects are pushed toward the loadings as well (as values can vary

more, thus potentially weakening the relationship between observed and latent variable).

An example of interpretation on real data is given in a later section. From a research

study, only one effect size for each equality constraint would be calculated. The

interpretation will often be up to the researcher's smallest effect of interest, and this

simulation gives some guidance that the values should not be interpreted with traditional

rules of thumb. The pattern of effects is potentially the most useful information: 1) positive

effects on the loadings with negative or very close to zero effects on the other parameters

may indicate a non-replication in loadings, 2) equal effects on loadings and intercepts with

Figure B5

Medium Intercepts Model Visualization

smaller or negative effects may indicate intercepts may be an issue, and 3) residuals may be
determined by the same pattern as loadings but with a smaller ratio of loadings to residuals
effect (i.e., loadings h_{nmi} / residuals h_{nmi}. The "size" could be determined by the ratio of
effect sizes for each constraint. Of course, this represents one simulation study, and results
from many studies in a meta-analysis would be fruitful for future work.

Parameter Replication and Effect Size

Figure B11 shows the effect size differences within large loadings simulations. The
results demonstrate that most of the loadings were considered non-invariant in the
bootstrapped models (while holdings all others equal). This result is partially due to
simulating very good data, so small changes in loadings results in a drop in fit for our chosen

Figure B6

Large Intercepts Model Visualization

invariance index. However, we can use this graph to show that question four shows a

possible effect size ranging from -0.07 to 0.13. The h_{nmi_p} value for question four was 0.27,

representing about a quarter of a possible total effect. Last, the density plot in Figure B12

shows the separation of the two different groups loadings in item four, thus, illustrating

group differences in the findings for their loadings. Each of the other combination of plots

can be found in the supplemental materials.

MGCFA Model Fit Statistics

Model fit statistics are provided for each of the ten model combinations (invariant,

three sizes for each loadings, intercepts, and residuals). These tables could be used to

examine the traditional change in fit statistics cutoff rules of thumb (Cheung & Rensvold,

Item Invariance

Latent Mean Distribution

Figure B7

Small Residuals Model Visualization

1161 2002), such as Δ CFI or Δ RMSEA, to the visualizations presented in the manuscript.

Figure B8

Medium Residuals Model Visualization

Table B1

Model Fit for Invariant Model

Model	AIC	BIC	CFI	TLI	RMSEA	SRMR
Overall	7,516.454	7,579.673	1.000	1.036	0.000	0.006
Group Group 1	3,765.749	3,818.571	0.976	0.953	0.047	0.031
Group Group 2	3,767.599	3,820.421	1.000	1.008	0.000	0.021
Configural	7,533.348	7,659.786	0.991	0.982	0.030	0.026
Loadings	7,528.476	7,638.056	0.994	0.992	0.020	0.033
Intercepts	7,522.397	7,615.118	1.000	1.003	0.000	0.035
Residuals	7,520.435	7,592.083	0.991	0.992	0.020	0.046

Figure B9

Large Residuals Model Visualization

Table B2

Model Fit for Small Differences in Loadings

Model	AIC	BIC	CFI	TLI	RMSEA	SRMR
Overall	7,530.321	7,593.540	0.977	0.955	0.049	0.025
Group Group 1	3,765.749	3,818.571	0.976	0.953	0.047	0.031
Group Group 2	3,785.242	3,838.064	0.979	0.958	0.050	0.029
Configural	7,550.991	7,677.430	0.978	0.956	0.048	0.030
Loadings	7,550.133	7,659.713	0.966	0.952	0.051	0.047
Intercepts	7,542.675	7,635.397	0.979	0.977	0.035	0.047
Residuals	7,534.091	7,605.739	0.993	0.994	0.019	0.054

Table B3

Model Fit for Medium Differences in Loadings

Model	AIC	BIC	CFI	TLI	RMSEA	SRMR
Overall	7,595.822	7,659.041	1.000	1.017	0.000	0.012
Group Group 1	3,765.749	3,818.571	0.976	0.953	0.047	0.031
Group Group 2	3,815.314	3,868.136	0.996	0.991	0.025	0.023
Configural	7,581.063	7,707.501	0.988	0.976	0.038	0.027
Loadings	7,609.348	7,718.928	0.878	0.826	0.101	0.079
Intercepts	7,601.550	7,694.271	0.891	0.879	0.084	0.079
Residuals	7,596.811	7,668.459	0.890	0.905	0.075	0.096

Table B4

Model Fit for Large Differences in Loadings

Model	AIC	BIC	CFI	TLI	RMSEA	SRMR
Overall	7,665.760	7,728.979	1.000	1.022	0.000	0.010
Group Group 1	3,765.749	3,818.571	0.976	0.953	0.047	0.031
Group Group 2	3,855.953	3,908.775	0.997	0.995	0.020	0.023
Configural	7,621.702	7,748.140	0.989	0.978	0.036	0.027
Loadings	7,659.690	7,769.270	0.852	0.789	0.114	0.088
Intercepts	7,652.360	7,745.081	0.863	0.848	0.097	0.088
Residuals	7,664.853	7,736.502	0.806	0.832	0.102	0.135

Table B5

Model Fit for Small Differences in Intercepts

Model	AIC	BIC	CFI	TLI	RMSEA	SRMR
Overall	7,520.471	7,583.690	1.000	1.035	0.000	0.007
Group Group 1	3,765.749	3,818.571	0.976	0.953	0.047	0.031
Group Group 2	3,767.599	3,820.421	1.000	1.008	0.000	0.021
Configural	7,533.348	7,659.786	0.991	0.982	0.030	0.026
Loadings	7,528.476	7,638.056	0.994	0.992	0.020	0.033
Intercepts	7,526.312	7,619.034	0.987	0.986	0.027	0.040
Residuals	7,524.356	7,596.005	0.975	0.978	0.033	0.050

Table B6

Model Fit for Medium Differences in Intercepts

Model	AIC	BIC	CFI	TLI	RMSEA	SRMR
Overall	7,538.375	7,601.594	1.000	1.033	0.000	0.008
Group Group 1	3,765.749	3,818.571	0.976	0.953	0.047	0.031
Group Group 2	3,767.599	3,820.421	1.000	1.008	0.000	0.021
Configural	7,533.348	7,659.786	0.991	0.982	0.030	0.026
Loadings	7,528.476	7,638.056	0.994	0.992	0.020	0.033
Intercepts	7,544.002	7,636.724	0.917	0.907	0.068	0.059
Residuals	7,542.064	7,613.712	0.905	0.917	0.065	0.067

Table B7

Model Fit for Large Differences in Intercepts

Model	AIC	BIC	CFI	TLI	RMSEA	SRMR
Overall	7,568.748	7,631.967	1.000	1.032	0.000	0.008
Group Group 1	3,765.749	3,818.571	0.976	0.953	0.047	0.031
Group Group 2	3,767.599	3,820.421	1.000	1.008	0.000	0.021
Configural	7,533.348	7,659.786	0.991	0.982	0.030	0.026
Loadings	7,528.476	7,638.056	0.994	0.992	0.020	0.033
Intercepts	7,574.054	7,666.776	0.797	0.775	0.106	0.084
Residuals	7,572.174	7,643.823	0.785	0.813	0.097	0.090

Table B8

Model Fit for Small Differences in Residuals

Model	AIC	BIC	CFI	TLI	RMSEA	SRMR
Overall	7,462.007	7,525.226	1.000	1.020	0.000	0.013
Group Group 1	3,765.749	3,818.571	0.976	0.953	0.047	0.031
Group Group 2	3,703.797	3,756.619	0.962	0.924	0.061	0.037
Configural	7,469.546	7,595.984	0.969	0.938	0.054	0.034
Loadings	7,471.637	7,581.217	0.944	0.920	0.062	0.049
Intercepts	7,465.722	7,558.443	0.952	0.946	0.051	0.051
Residuals	7,465.986	7,537.635	0.930	0.939	0.054	0.065

Table B9

Model Fit for Medium Differences in Residuals

Model	AIC	BIC	CFI	TLI	RMSEA	SRMR
Overall	7,382.876	7,446.095	0.992	0.985	0.028	0.020
Group Group 1	3,765.749	3,818.571	0.976	0.953	0.047	0.031
Group Group 2	3,601.920	3,654.742	1.000	1.023	0.000	0.018
Configural	7,367.669	7,494.108	0.996	0.992	0.020	0.025
Loadings	7,371.027	7,480.607	0.969	0.956	0.049	0.046
Intercepts	7,364.724	7,457.446	0.977	0.975	0.037	0.047
Residuals	7,386.854	7,458.502	0.877	0.893	0.076	0.076

Table B10

Model Fit for Large Differences in Residuals

Model	AIC	BIC	CFI	TLI	RMSEA	SRMR
Overall	7,301.897	7,365.116	0.993	0.986	0.026	0.019
Group Group 1	3,765.749	3,818.571	0.976	0.953	0.047	0.031
Group Group 2	3,454.309	3,507.131	0.945	0.889	0.075	0.036
Configural	7,220.058	7,346.496	0.960	0.920	0.063	0.034
Loadings	7,218.878	7,328.458	0.948	0.926	0.060	0.046
Intercepts	7,213.011	7,305.732	0.956	0.951	0.049	0.048
Residuals	7,305.660	7,377.309	0.559	0.617	0.137	0.189

Figure B10

Visualization of the effect size of bootstrapped replication proportions on simulated data. Each panel indicates the simulated data type, colors represent the differences in the strength of the non-invariance, and the bars on the x-axis represent the effect size for the equality constraint.

Figure B11

Bootstrapped and Random Group effect size differences in loadings for the Large Loading difference simulation. The size of the point reprensents the number of data points included in that calculation.

Figure B12

Bootstrapped and Random density plots for invariant and non-invariant bootstrapped partial effects examining only large loadings.

Table C1

Fit Estimates for Partial Invariance

Residuals on Invariant Data

Estimated Parameter	CFI	RSMEA
q1 ~~ q1	0.990	0.021
q2 ~~ q2	0.987	0.024
q3 ~~ q3	0.996	0.014
q4 ~~ q4	1.000	0.000
q5 ~~ q5	0.987	0.025
lv ~~ lv	0.991	0.020

Appendix C

Simulated Partial Invariance Results

Table C2

*Fit Estimates for Partial Invariance
Loadings for Small Loading Data*

Estimated Parameter	CFI	RSMEA
lv =~ q1	0.993	0.019
lv =~ q2	0.989	0.023
lv =~ q3	0.989	0.023
lv =~ q4	1.000	0.000
lv =~ q5	0.994	0.017

Table C3

*Fit Estimates for Partial Invariance
Loadings for Medium Loading Data*

Estimated Parameter	CFI	RSMEA
lv =~ q1	0.890	0.075
lv =~ q2	0.904	0.072
lv =~ q3	0.887	0.078
lv =~ q4	1.000	0.000
lv =~ q5	0.914	0.068

Table C4

Fit Estimates for Partial Invariance

Loadings for Large Loading Data

Estimated Parameter	CFI	RSMEA
lv =~ q1	0.806	0.102
lv =~ q2	0.812	0.102
lv =~ q3	0.813	0.102
lv =~ q4	1.000	0.000
lv =~ q5	0.861	0.088

Table C5

Fit Estimates for Partial Invariance

Loadings for Small Intercept Data

Estimated Parameter	CFI	RSMEA
q1 ~1	0.975	0.033
lv ~1	0.975	0.033
q2 ~1	0.972	0.035
q3 ~1	0.972	0.036
q4 ~1	0.988	0.023
q5 ~1	0.971	0.036

Table C6

Fit Estimates for Partial Invariance

Loadings for Medium Intercept Data

Estimated Parameter	CFI	RSMEA
q1 ~1	0.905	0.065
lv ~1	0.905	0.065
q2 ~1	0.901	0.067
q3 ~1	0.901	0.067
q4 ~1	0.988	0.023
q5 ~1	0.902	0.067

Table C7

Fit Estimates for Partial Invariance

Loadings for Large Intercept Data

Estimated Parameter	CFI	RSMEA
q1 ~1	0.785	0.097
lv ~1	0.785	0.097
q2 ~1	0.781	0.100
q3 ~1	0.781	0.100
q4 ~1	0.988	0.023
q5 ~1	0.784	0.099

Table C8

Fit Estimates for Partial Invariance

Loadings for Small Residual Data

Estimated Parameter	CFI	RSMEA
q1 ~~ q1	0.928	0.056
q2 ~~ q2	0.936	0.053
q3 ~~ q3	0.926	0.057
q4 ~~ q4	0.955	0.044
q5 ~~ q5	0.926	0.057
lv ~~ lv	0.930	0.054

Table C9

Fit Estimates for Partial Invariance

Loadings for Medium Residual Data

Estimated Parameter	CFI	RSMEA
q1 ~~ q1	0.879	0.077
q2 ~~ q2	0.873	0.079
q3 ~~ q3	0.878	0.077
q4 ~~ q4	0.980	0.031
q5 ~~ q5	0.873	0.079
lv ~~ lv	0.877	0.076

Table C10

Fit Estimates for Partial Invariance

Loadings for Large Residual Data

| Estimated Parameter	CFI	RSMEA
q1 ~~ q1	0.555	0.141
q2 ~~ q2	0.556	0.141
q3 ~~ q3	0.556	0.141
q4 ~~ q4	0.967	0.039
q5 ~~ q5	0.557	0.141
lv ~~ lv	0.559	0.137

Appendix D

Invariance Plots Difference Scores by Condition

Invariant Graph Results

Small Loadings Graph Results

Medium Loadings Graph Results

Large Loadings Graph Results

Small Intercepts Graph Results

Medium Intercepts Graph Results

Large Intercepts Graph Results

Small Residuals Graph Results

Medium Residuals Graph Results

Large Residuals Graph Results

Appendix E

Invariance Plots Effect Sizes by Condition

Invariant Graph Results

Small Loadings Graph Results

Medium Loadings Graph Results

Large Loadings Graph Results

Small Intercepts Graph Results

Medium Intercepts Graph Results

Large Intercepts Graph Results

Small Residuals Graph Results

Medium Residuals Graph Results

Large Residuals Graph Results

Appendix F

Density Plots by Condition

Invariant Graph Results

Small Loadings Graph Results

Medium Loadings Graph Results

Large Loadings Graph Results

Small Intercepts Graph Results

Medium Intercepts Graph Results

Large Intercepts Graph Results

Small Residuals Graph Results

Medium Residuals Graph Results

Large Residuals Graph Results

Replication Test

Data

```
## # A tibble: 2 x 2
##   group       sample
##   <chr>        <int>
## 1 Chen          1010
## 2 Schulenberg   1765
```

MGCFA

```
## # A tibble: 7 x 7
##   model              AIC      BIC   cfi   tli  rmsea   srmr
##   <chr>            <dbl>    <dbl> <dbl> <dbl>  <dbl>  <dbl>
## 1 Overall        113442. 113608. 0.939 0.928 0.0890 0.0344
## 2 Group Schulenberg 69101.  69254. 0.928 0.915 0.108  0.0353
```

```
## 3 Group Chen            42603.   42741. 0.929 0.916 0.0766 0.0399

## 4 Configural          111760. 112258. 0.929 0.916 0.0975 0.0347

## 5 loadings            111789. 112210. 0.927 0.921 0.0946 0.0427

## 6 intercepts          112785. 113129. 0.892 0.891 0.111  0.0599

## 7 residuals           113318. 113579. 0.873 0.880 0.116  0.0661
```

Overall, the one-factor model fits the data well. Each group also shows adequate model fit. If we use ΔCFI $<= .01$, we find that the loadings would be considered invariant across the English and Chinese samples. The intercepts were not invariant.

Partial Invariance

```
## # A tibble: 15 x 2

##    free.parameter cfi

##    <chr>          <lvn.vctr>

##  1 "RS2 ~1 "      0.8970997

##  2 "RS11 ~1 "     0.8963948

##  3 "RS4 ~1 "      0.8961938

##  4 "RS12 ~1 "     0.8960095

##  5 "RS14 ~1 "     0.8954496

##  6 "RS1 ~1 "      0.8954215

##  7 "RS10 ~1 "     0.8953395

##  8 "RS3 ~1 "      0.8948455

##  9 "RS5 ~1 "      0.8945236

## 10 "RS7 ~1 "      0.8941450

## 11 "RS8 ~1 "      0.8920022

## 12 "RS13 ~1 "     0.8919491

## 13 "RS9 ~1 "      0.8919293

## 14 "RS ~1 "       0.8917757
```

```
## 15 "RS6 ~1 "          0.8917604
```

Examining partial invariance reveals several potential candidates for partial invariance. In this next section, we relaxed group constraints until we achieved partial invariance (i.e., ΔCFI $<=$ .01). We will need to find our CFI as at least 0.92. More than half the items are necessary to achieve "partial" invariance (which really implies no invariance is likely possible).

```
## # A tibble: 1 x 6
##        AIC      BIC   cfi   tli  rmsea   srmr
##      <dbl>    <dbl> <dbl> <dbl>  <dbl>  <dbl>
## 1 111935. 112326. 0.922 0.917 0.0966 0.0462
```

```
## # A tibble: 6 x 3
##   term        English Chinese
##   <chr>         <dbl>   <dbl>
## 1 "RS1 ~1 "      5.24    4.78
## 2 "RS3 ~1 "      5.14    5.46
## 3 "RS4 ~1 "      5.29    5.71
## 4 "RS5 ~1 "      5.13    5.13
## 5 "RS7 ~1 "      5.26    5.26
## 6 "RS12 ~1 "     5.57    5.16
```

```
##  RS1  RS2  RS3  RS4  RS5  RS6  RS7  RS8  RS9 RS10 RS11 RS12 RS13 RS14
## 0.32 0.29 0.23 0.30 0.00 0.00 0.00 0.00 0.00 0.29 0.27 0.29 0.00 0.26
```

Visualize Invariance

```
## Plot for RS2
```

Item Invariance

## Latent Mean Distribution

Group — Schulenberg — Chen

1255 ## -------

1256 ## Plot for RS11

1258 ## -------

1259 ## Plot for RS4

Item Invariance

Latent Mean Distribution

Group —✕— Schulenberg —✕— Chen

1260

1261 ## -------

1262 ## Plot for RS12

Item Invariance

Latent Mean Distribution

Group — Schulenberg — Chen

1263

1264   ## -------

1265   ## Plot for RS14

1267 ## -------

1268 ## Plot for RS1

Item Invariance

Latent Mean Distribution

Group — Schulenberg — Chen

1269

1270   ## -------

1271   ## Plot for RS10

1273 ## -------

1274 ## Plot for RS3

Bootstrap Model

```
##        model non_invariant random_non_invariant     h_nmi h_nmi_p
## 1 intercepts              1                    0 3.141593       1
```

In this case, we do not see loadings print out. That implies that all models in both real data and randomized data are invariant because the function only calculates information for non-invariance. We see that the intercepts are unlikely to ever replicate across Chinese and English samples. This result is not surprising given the large number of relaxed parameters required to achieve partial invariance.

Bootstrap Partial Invariance

Each Parameter on the Overall Model Invariance

```
##        term non_invariant random_non_invariant     h_nmi  h_nmi_p
```

```
## 1     RS ~1              1              0.002 3.052120 0.971520
## 2    RS1 ~1              1              0.001 3.078337 0.979865
## 3   RS10 ~1              1              0.002 3.052120 0.971520
## 4   RS11 ~1              1              0.002 3.052120 0.971520
## 5   RS12 ~1              1              0.002 3.052120 0.971520
## 6   RS13 ~1              1              0.002 3.052120 0.971520
## 7   RS14 ~1              1              0.002 3.052120 0.971520
## 8    RS2 ~1              1              0.002 3.052120 0.971520
## 9    RS3 ~1              1              0.002 3.052120 0.971520
## 10   RS4 ~1              1              0.002 3.052120 0.971520
## 11   RS5 ~1              1              0.002 3.052120 0.971520
## 12   RS6 ~1              1              0.002 3.052120 0.971520
## 13   RS7 ~1              1              0.001 3.078337 0.979865
## 14   RS8 ~1              1              0.002 3.052120 0.971520
## 15   RS9 ~1              1              0.002 3.052120 0.971520
```

In this output, we see that all the bootstrapped runs of the real data are non-invariant, even when each parameter is relaxed individually. A few runs of the random data are *non*-invariant (meaning most are actually invariant when randomized). This indicates that no one parameter is likely the reason for non-invariance, as they all show large non-replication effects. If we use the `boot_summary`, we can see the effect size for each parameter when the two intercepts are compared to each other (as the chart above shows the overall model invariance effect).

Each Parameter's Standardized Difference Score

```
##         term invariant n_boot     d_boot
## 1     RS ~1     FALSE   1000         NA
## 2    RS1 ~1     FALSE   1000  0.2913872
```

```
## 3   RS10 ~1     FALSE   1000  0.3596961

## 4   RS11 ~1     FALSE   1000 -0.6589317

## 5   RS12 ~1     FALSE   1000 -0.3010202

## 6   RS13 ~1     FALSE   1000 -0.4005088

## 7   RS14 ~1     FALSE   1000 -0.5665988

## 8    RS2 ~1     FALSE   1000 -0.5540349

## 9    RS3 ~1     FALSE   1000 -0.6218628

## 10   RS4 ~1     FALSE   1000 -0.6245830

## 11   RS5 ~1     FALSE   1000 -0.3334520

## 12   RS6 ~1     FALSE   1000 -0.5014247

## 13   RS7 ~1     FALSE   1000 -0.2381644

## 14   RS8 ~1     FALSE   1000 -0.6333750

## 15   RS9 ~1     FALSE   1000 -0.5986866


##          term invariant n_random d_random

## 1     RS ~1     FALSE          2       NA

## 2    RS1 ~1     FALSE          1       NA

## 3   RS10 ~1     FALSE          2       NA

## 4   RS11 ~1     FALSE          2       NA

## 5   RS12 ~1     FALSE          2       NA

## 6   RS13 ~1     FALSE          2       NA

## 7   RS14 ~1     FALSE          2       NA

## 8    RS2 ~1     FALSE          2       NA

## 9    RS3 ~1     FALSE          2       NA

## 10   RS4 ~1     FALSE          2       NA

## 11   RS5 ~1     FALSE          2       NA

## 12   RS6 ~1     FALSE          2       NA

## 13   RS7 ~1     FALSE          1       NA
```

```
## 14  RS8 ~1        FALSE         2         NA

## 15  RS9 ~1        FALSE         2         NA
```

With the bootstrap summary, we see that no invariant intercept runs are found (not all columns shown to save space). Note that the `RS~1` does not calculate d because the sd is zero (thus, this would cause an error as sd is part of the denominator). The random runs also do not show a d score because the effect size is only calculated when at least 10% of the runs in `n_boot` or `n_random` are found. Since we only have a few in that column, no effect size is calculated. The differences in the intercepts for the real data (`boot` columns) appear to be medium to large, showing d scores from 0.3 to 0.6. Therefore, we might expect that the English and Chinese samples have different average endorsement levels of the RS14.
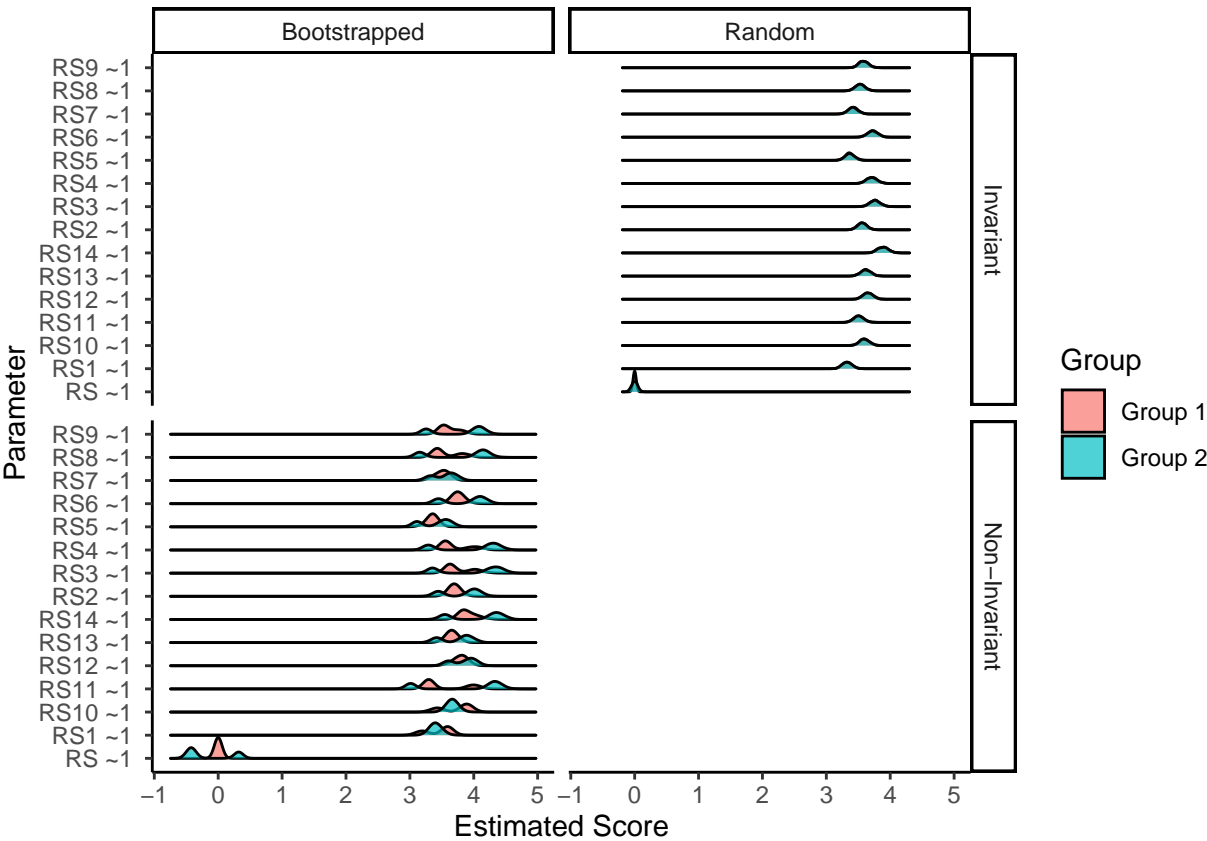

The examination of the density plot shows how group 1 (Schulenberg English) tends to show lower average scores than group 2 (Chen Chinese) for most but not all items.

1354   We can view the mean difference or standardized mean difference by using:

When effects are non-invariant in the randomized data, the mean difference is still fairly small, but we see large mean differences in intercepts when the bootstrapped data is non-invariant. In the effect size graph, we can see that this effect is medium to large for all the parameters.